

Political Violence and Bureaucratic Discretion: Evidence from German Asylum Decisions

Yogam Tchokni*

03/07/2026

[\[Click here for the most recent version\]](#)

Abstract

Democratic accountability typically flows downward: voters pressure politicians, who direct agencies. I test whether policy also responds from below—through the discretionary decisions of the bureaucrats who implement it. I study German asylum decisions, exploiting arson attacks against refugee accommodations as salient signals of anti-immigrant hostility. Unlike out-group-perpetrated violence studied previously, in-group attacks against a vulnerable minority activate competing pressures on decision-makers: compliance with hostile demands versus empathy for victims. Using administrative data on 95,945 applicants (2015–2018), immigration offices exposed to a nearby attack show a 0.6 percentage-point decline in the monthly approval probability (23 percent relative to the mean), without a change in policy directives. Survey evidence ties the decline to the interview month; suggestive evidence points to larger effects for attacks receiving more news coverage.

Keywords: Immigration, Political Violence, Asylum, Discrimination, Bureaucratic Discretion

JEL codes: D73, J15, K37

*DIW Berlin, ytchokni@diw.de. I thank Daniel Graeber, Lorenz Meister, Neil Murray, and Laura Steinkemper for sharing the raw LexisNexis data.

1 Introduction

What if bureaucrats take the law into their own hands? Democratic accountability is typically understood as flowing downward: voters pressure politicians, who direct agencies. But policy implementation depends on the people who carry it out. Across courts, welfare offices, and government agencies, public servants exercise substantial discretion (Lipsky, 2010). When implementers possess such power, their decisions may respond not only to demands from policymakers but also to the social and political conditions of the local environment in which they operate. Policy may therefore be responsive to demands from the *bottom-up* through the individual’s use of discretionary power beyond traditional *top-down* changes in directives by politicians.

To test whether signals of extreme and violent opposition to government policy affect its execution through public servants’ use of bureaucratic discretion, I use the setting of German asylum decisions and the wave of political violence against refugees from 2015 to 2018. I show that *in-group* violence through arson attacks against refugee housing directly reduces asylum approval rates, through the personal discretion of immigration officers.

The German asylum system is well-suited to study this question. First, asylum policy is determined at the federal level, so local variation in approval rates reflects adjudication rather than policy change. Second, applicants are quasi-randomly allocated across Germany through a centralized distribution system, limiting sorting. Third, geographic separation between applicants’ residences and the 48 branch offices that process their claims allows me to separate office-side exposure to local hostility from exposure near applicants’ own residences. Arson attacks against refugee accommodations, perpetrated by members of the majority population against a vulnerable minority, serve as salient, localized signals of extreme and violent opposition to existing government policy.

I find consistent evidence across three independent data sources that arson attacks against refugee accommodations reduce asylum approval rates. The main analysis uses individual-level administrative data from the Central Register of Foreigners (95,945 applicants, 2015–2018) and shows that immigration offices exposed to a nearby attack exhibit a 0.6 percentage-point decline in the monthly approval probability (23 percent relative to the mean), without a change in policy directives. A within-office test comparing cases decided via written procedure to those decided via a personal interview at the same office during the same period confirms that the effect operates through the interview: caseworkers who meet applicants face-to-face adjust their decisions after nearby attacks, while those deciding cases on paper do not. A second, independently collected dataset of office-level aggregate statistics on first-instance decisions corroborates the decline and shows that attacks shifted the content of decisions rather than their timing or volume. A third test uses individual-level survey data from the IAB-BAMF-SOEP refugee survey; the decline in approval probability replicates and is concentrated at the interview month, pinpointing

the moment where personal discretion of immigration officers is highest as the moment where local conditions enter adjudication. Results are consistent with a *bottom-up* channel in which approval rates comply with the demands of political violence.

In the refugee survey, the decline is larger for attacks receiving substantial media coverage, consistent with an information channel identified in previous research (Emeriau, 2024; Spirig, 2023). Survey evidence from the SOEP shows that arson attacks shift public attitudes among residents of immigration office counties. Concerns about immigration rise persistently, paralleling the durable voting shifts documented in Sabet et al. (2023). Concerns about xenophobic hostility spike in the immediate aftermath but dissipate within weeks, making any resistive responses short-lived, consistent with previous results (Freitas-Monteiro and Prömel, 2024). The asymmetry provides attitudinal evidence for why the bureaucratic response tilts toward stricter treatment.

A large body of literature documents how political violence reshapes political outcomes, tracing effects almost exclusively through the electoral channel. Terrorist attacks shift voting toward nationalist parties (Gould and Klor, 2010; Aidt and Franck, 2015; Rees and Smith, 2022; Gassebner et al., 2008), and far-right violence increases support for far-right parties and hardens anti-immigration attitudes (Sabet et al., 2023; Jomni and Zakharov, 2026), though not universally (Larsen et al., 2020). In each case, violence enters the policy process by changing how citizens vote, which pressures politicians to adjust.

Bureaucratic discretion shapes outcomes across government: in courts (Shayo and Zussman, 2011; Arnold et al., 2018; Eren and Mocan, 2018; Dobbie et al., 2018), policing (Grosjean et al., 2023), welfare allocation (Hemker and Rink, 2017; Maestas et al., 2013; Bolhaar et al., 2020), and immigration (Emeriau, 2023, 2024; Chen et al., 2016; Gundacker et al., 2024). External events influence how that discretion is exercised (Gould and Klor, 2016; Chen and Loecher, 2025; Heyes and Saberian, 2019; Danziger et al., 2011). In the literature closest to this paper, Shayo and Zussman (2011) show that Israeli judges exhibit heightened in-group bias after terrorist attacks, and Brodeur and Wright (2019) find that 9/11 reduced asylum approval rates for Muslim applicants in the United States—both cases of out-group-perpetrated violence where stricter treatment admits a rational-threat-updating interpretation. Emeriau (2024) shows that news of Mediterranean shipwrecks raised asylum approval rates in France, demonstrating that empathy can push in the opposite direction. The most closely related studies are Spirig (2023), who shows that media coverage of immigration reduces judicial approval rates in Switzerland, and Gundacker et al. (2024), who find lower approval probabilities for refugees in immigration-averse regions of Germany. Neither examines whether political violence directly shapes bureaucratic decisions.

This paper makes two contributions. First, rather than tracing responses to *in-group* violence through the ballot box, I examine whether it alters policy implementation directly,

through the discretionary decisions of the public servants who carry it out. Compared to other studies of violent events and adjudication (Emeriau, 2024; Brodeur and Wright, 2019; Shayo and Zussman, 2011), I provide evidence on a qualitatively distinct hostility signal through the *in-group*. Unlike out-group-perpetrated attacks, the predicted direction of the bureaucratic response is ambiguous. Such violence could normalize hostility (Bursztyn et al., 2020; Sabet et al., 2023) and serve as a signal of extreme opposition to current policy, demanding stricter bureaucratic treatment. Violent or extremist acts may also discredit political positions (Bursztyn et al., 2020; Chenoweth and Stephan, 2011; Feinberg et al., 2020; Freitas-Monteiro and Prömel, 2024) and create empathy for the targeted group. Existing work on out-group-perpetrated violence that results in stricter treatment can be interpreted through rational decision-making given updated threat perceptions. In-group violence against a vulnerable out-group offers no such mechanism, as the targeted group poses no greater threat after the attack. Instead, it generates competing pressures on decision-makers: compliance with hostile demands versus empathy-driven leniency. I show that the former wins out.

Second, I provide the first analysis of asylum adjudication using individual-level administrative data from Germany, the largest recipient of asylum applications in Europe. Between 2015 and 2018, Germany received approximately 1.5 million first-time applications, accounting for 40 percent of all applications filed across Europe. Previous studies using administrative data have focused on countries with substantially smaller caseloads such as France (9 percent) (Emeriau, 2024) and Switzerland (3 percent) (Spirig, 2023).

Section 2 describes the German asylum system. Section 3 develops the conceptual framework. Section 4 presents the data, Section 5 the empirical strategy, and Section 6 the results. Section 8 concludes.

2 Institutional Setting

2.1 Allocation System

Germany administers refugee arrivals through a centralized allocation system. Upon arrival, asylum seekers register at reception centers and are allocated to federal states via the EASY system (*Erstverteilung der Asylbegehrenden*), which follows the Königsteiner Schlüssel, a formula weighted by each state’s population and tax revenue. States then distribute individuals to municipalities based on housing availability. Allocation decisions are largely independent of individual refugee characteristics (Jaschke, Sardoschau, and Tabellini, 2023; Aksoy et al., 2023).

Following municipal assignment, applicants are directed to one of 48 branch offices for processing, often located in different counties than applicants’ residences. This geographic wedge is what allows me to separate office-side exposure from applicant-side exposure:

an arson attack near an immigration office reaches that office’s staff and local environment. Given that applicants in the same location may be assigned to different offices, confounding effects on applicant behavior can be controlled for.

2.2 Decision Process

Each applicant undergoes a personal interview conducted by an asylum officer. The interview forms the basis for assessing eligibility for protection. Officers have substantial discretion: they control the line of questioning, assess the credibility of testimony, and ultimately determine both whether to grant protection and what level of protection to award.

Applicants may receive one of several outcomes: full refugee status under the Geneva Convention (three-year permit, immediate family reunification), subsidiary protection (one-year permit, restricted reunification), a deportation ban, or rejection. Rejected applicants may appeal to administrative courts. See Appendix Table A.1 for detailed descriptions of protection levels. Figure 2 illustrates the observability of each outcome type under the 12-month restriction used in the empirical design.

The interview serves as the primary opportunity for decision-makers to assess the credibility and substance of an asylum claim. During this meeting, officers evaluate the applicant’s testimony, probe for inconsistencies, and form judgments about protection eligibility. While applications are filed and administrative processing occurs over extended periods, the interview represents the concentrated moment when subjective evaluation takes place. Events occurring close to this critical decision point are therefore most likely to influence officer judgments, whether through heightened threat perceptions following local violence or through increased empathy following humanitarian tragedies. Following Emeriau (2024), who uses interview dates to identify the timing of discretionary effects on asylum decisions in France, I treat the interview as the critical moment for adjudication.

The degree of discretion varies by applicant origin. During 2016–2018, Syrian applicants were rarely rejected outright, though officers exercised discretion over whether to grant full refugee status or the less favorable subsidiary protection.¹ For applicants from other origins, officer discretion extends to the binary approval decision as well. This variation in the scope of discretion generates a testable prediction: if attacks operate through officer discretion at the interview, effects should concentrate among nationalities where officers have more latitude over outcomes.

The scope for discretion was particularly wide at the start of the study period due to weak internal oversight. The immigration office’s central quality-assurance unit reviewed

¹In 2016, 167 of 266,866 Syrian applications were rejected outright; of the remainder, 165,754 received full Geneva-Convention protection and the balance received subsidiary protection, which restricts family reunification.

approximately 1% of decisions in 2015, and the dual-review requirement for individual decisions was made formally binding by central directive only in September 2017 (see Appendix A for details). Until at least late 2017, officers operated with minimal systematic review of their individual decisions.

3 Conceptual Framework

I conceptualize arson attacks against refugee accommodations as salient events that increase the local visibility of anti-immigrant sentiment and social dissatisfaction. Such attacks signal that segments of the population hold strong anti-refugee views, intense enough to motivate violent action. The existing literature on how salient events affect bureaucratic and judicial decision-makers suggests several pathways through which such attacks may alter officer behavior. The predicted direction is ambiguous *ex ante*, as the same event can produce both compliance with and resistance to political violence.

Shift in own beliefs. Arson attacks signal that social cohesion in the local community is under strain. Officers who learn of violence against refugee accommodations may perceive their surroundings as destabilized. This perceived disruption can harden attitudes toward the group associated with the tension, even when that group is itself the target. Sabet et al. (2023) show that successful right-wing attacks in Germany shift individual attitudes rightward and increase anti-immigration concern, even among previously unaffiliated individuals. Salient events do not need to change beliefs from nothing but can instead activate predispositions that already exist. Grosjean et al. (2023) show that Trump campaign rallies increased racial bias in policing, with effects concentrated among officers whose baseline behavior was already more biased. In the immigration office context, rapid hiring during 2015–2016 brought officers with heterogeneous predispositions into the same offices. An arson attack nearby may both shift attitudes on the margin and activate restrictive tendencies among officers already predisposed—two sides of the same process, requiring neither strategic calculation nor conscious awareness.

Shift in perceived expectations. A distinct but complementary channel operates through perceived expectations rather than personal beliefs. Attacks may not change what officers think about refugees, but instead what officers believe their environment expects of them. Spirig (2023) shows that higher asylum salience in Swiss media leads to more restrictive judicial decisions on asylum appeals and is not restricted to judges of any particular ideology. Officers may calibrate decisions toward what they perceive as the locally appropriate position, whether through internalized salience effects or through social pressure from colleagues and the community atmosphere (Lipsky, 2010). This channel does not require officers to personally hold anti-refugee views. The institutional context of 2016–2017 amplifies this pathway: the immigration office was rapidly expanding its

workforce, and new hires on probationary periods may have been especially susceptible to signals about local expectations. Emeriau (2023) shows that new hires are often more biased than experienced ones.

Protective action. Arson attacks physically endanger refugees in their accommodations. Given that positive asylum decisions allow refugees to choose their place of residence more freely, officers confronted with physical endangerment may view applicants more favorably. Emeriau (2024) shows that news of Mediterranean shipwrecks increased asylum approval rates in France by 4.4 percentage points when featured on prime-time television. However, this empathy effect was extremely short-lived, disappearing within one day of the event. Given that arson attacks occur within a context of social tension rather than humanitarian disaster, it is unclear whether they trigger the same victim framing.

Backlash against extremism. A broader literature on responses to political violence shows that extreme tactics frequently activate moral outrage and sympathy for targets (Wasow, 2020; Feinberg et al., 2020; Chenoweth and Stephan, 2011). In the German context, Freitas-Monteiro and Prömel (2024) document how far-right demonstrations increase concern about hostility toward foreigners, raise intentions to donate to refugee aid, and benefit left-wing parties. Officers may feel revulsion toward the attackers and compensate by exercising discretion more favorably.

The framework does not provide a clear prediction for the sign of the net effect. The resistance pathways require specific conditions, such as victim framing and immediate temporal proximity, that may not hold for arson attacks in this context. Attack severity provides one dimension of leverage: if empathy or backlash operates, aggravated arson that endangers human life should show a different pattern than minor property damage. Whether attacks receive media coverage may matter for the salience channel. I return to the severity distinction in the heterogeneity results (Appendix C.1) and to the media-coverage channel in the survey-based analysis (Section 7.2).

4 Data and Measurement

This section describes the data infrastructure used to analyze how localized political violence impacts asylum adjudication: official administrative records on anti-refugee violence compiled from parliamentary interpellations, individual-level longitudinal administrative tracking data from the Central Register of Foreigners, individual-level survey data from the IAB-BAMF-SOEP Survey of Refugees, office-level aggregate decision statistics, and news coverage data from LexisNexis.

4.1 Attack data

A key source of administrative data on anti-refugee violence comes from parliamentary interpellations submitted to the German federal government. These interpellations are formal inquiries that require the government to provide detailed information on specific topics, in this case regarding attacks on refugee accommodations and asylum seekers.

The federal government compiles responses using data from the Federal Criminal Police Office through the Criminal Police Reporting Service for Politically Motivated Crime. These reports provide incident-level data on attacks on asylum accommodations (both occupied facilities and those under construction), attacks on asylum seekers in public spaces, and far-right protests at or near refugee facilities. Each incident is recorded with the date, location (city and state), and the specific criminal offense.

I use quarterly interpellations covering 2015–2021, which provide the most systematic and comprehensive administrative record of anti-refugee violence during this period. The full 2015–2021 window is used to determine the timing of each office’s first nearby attack. Offices with no nearby attack during this extended window are classified as never-treated. Figure 1 displays the geographic distribution of these attacks across German counties over time.

The data have two advantages over media-based compilations. First, they are based on official police reports rather than media coverage, which reduces concerns about selective reporting. Second, they provide standardized legal classifications of offenses, enabling consistent severity coding across incidents and years.

The data classify arson incidents according to the German Criminal Code. I distinguish between simple arson and aggravated arson, where the latter involves endangering human life, occupied buildings, or serious bodily harm.² I construct an *aggravated arson* indicator for attacks classified under the aggravated categories, representing offenses with higher potential risk to human life. This distinction allows me to examine whether the severity of attacks moderates their effect on asylum decisions.

4.2 Application data

My main data source is the Central Register of Foreigners (CRF), a nationwide administrative register that is centrally managed by a federal office. It contains personal data, nationality, residency status, and other legal matters such as asylum applications, for all non-German citizens residing in or having resided in the country. The immigration office provides the AZR Research Dataset 2023, a 20% random sample drawn from the CRF on 31 December 2023 (Gleiser and Salas Escobar, 2024); for a detailed description, see the accompanying data report (Gleiser and Hinz, 2024).

²See Appendix Table A.2 for the full classification of criminal offenses.




Figure 1: Arson Attacks on Refugee Accommodations by County and Year

Notes: Data from parliamentary interpellations to the German federal government, compiled from Federal Criminal Police Office records. Counties are shaded by the number of arson attacks on refugee accommodations recorded in each year. The 2020–2023 panel aggregates the post-sample period for comparison.

The data contain information on place of residence, residence status, and other personal characteristics, with records entered by local municipalities, immigration offices, the police, and other authorities. For refugees, the residence-status history tracks the progression of the asylum application. Socio-demographic variables include age, gender, country of origin, education, and previous employment. Coverage is limited to non-citizens at the time of sample draw, since foreigners who naturalize are deleted from the CRF upon obtaining German citizenship. Appendix B.1 discusses whether selective sample attrition threatens identification.

The estimation sample covers person-month observations in calendar years 2015–2018. I remove applicants eligible for the expedited written procedure, in which decisions were rendered on the basis of a written questionnaire without a personal interview. Under this policy, Syrian, Iraqi, and Eritrean nationals who entered Germany before 1 January 2016 and filed their application before 17 March 2016 were decided through a fast-track pro-

cedure that bypassed the interview stage.³ Excluding these cases restricts the sample to applicants whose claims were adjudicated through the regular procedure with a personal interview, where officer discretion is exercised. Section 7.1 directly tests this mechanism by comparing decisions reached via personal interview to those decided through written procedure at the same office during the same period. The arson effect is concentrated entirely among interview cases.

The resulting sample restrictions and their effect on applicant counts are summarized in Table A.3. Treatment timing is computed over the full 2015–2021 interpellation window. Appendix Table A.4 lists each branch office, its treatment status, and first attack month. Table 1 breaks the estimation sample down by nationality and reports cumulative acceptance rates at different decision horizons. Syria and Eritrea account for roughly 20% of the sample and are resolved quickly (over 75% of positive decisions within 12 months), while Afghanistan and Iran have substantially lower and slower acceptance rates. Appendix Figure A.1 plots the discrete hazard rate and cumulative incidence of positive and negative decisions over the first 23 months after application.

Table 1: Acceptance Rates by Nationality and Decision Horizon

Nationality	<i>N</i>	Share (%)	≤3m (%)	≤6m (%)	≤12m (%)	≤24m (%)	Ever (%)
Other	44,856	46.8	1.9	4.0	7.0	10.2	13.8
Syria	16,821	17.5	26.4	51.6	75.5	88.3	93.8
Afghanistan	16,098	16.8	2.2	8.2	21.2	35.5	54.6
Iraq	8,559	8.9	11.0	26.8	49.0	55.0	59.2
Iran	4,086	4.3	2.9	9.8	27.6	44.5	59.8
Nigeria	2,393	2.5	0.0	0.5	3.4	7.3	12.5
Eritrea	2,103	2.2	23.9	42.2	64.4	76.7	81.3
Turkey	1,021	1.1	1.9	8.6	17.1	22.2	32.8
Total	95,937	100.0	7.5	16.1	27.3	35.1	42.3

Notes: Cumulative acceptance rates at different decision horizons for the estimation sample, pooled over application years 2015–2018. *N* is the number of unique applicants. Share is each nationality’s fraction of the total sample. Rates are cumulative: ≤6m includes all positive decisions rendered within 6 months of application, not only those between 3 and 6 months. “Ever” denotes the unconditional acceptance rate regardless of duration. Some year × nationality cells are suppressed under FDZ disclosure rules (counts below 5); suppressed cells are excluded from the numerator but not the denominator, so rates may slightly understate true values. Source: AZR Research Dataset 2023.

One key limitation of the CRF data is that it only contains legally binding changes to the asylum status. Therefore, asylum status is only observed once the legal appeals process has been exhausted (see Figure 2). Interview dates, initial decisions, and their timing therefore cannot be directly inferred from the data, so the treatment-relevant window cannot be precisely identified. However, the duration between application and final status


³From 2014 through late 2015, the immigration office granted refugee status to Syrian and Eritrean nationals and members of Iraqi religious minorities through a written fast-track procedure on the basis of a questionnaire alone. Beginning in 2016, personal interviews were reinstated for all applicants. See Section 7.1.

provides information on the initial decision and motivates the use of a duration model. The next section describes how these features are leveraged.

(a) Case A: Observed Approval



(b) Case B: Observed Rejection



(c) Case C: Appeals Outcome




Figure 2: Data Observability in Asylum Decision Processes

Notes: Schematic illustration of three types of asylum application outcomes and their observability under the 12-month restriction. Green shading indicates stages observed in the data; grey shading indicates stages that are unobserved. Case A: positive decision within 12 months (observed). Case B: negative decision within 12 months (observed as zero in outcome variable). Case C: appeal process extending beyond 12 months (applicant exits sample as censored). The 12-month window is designed to capture first-instance determinations while excluding most appeals.

4.3 Additional data

Survey data. I complement the administrative data with survey data from the IAB-BAMF-SOEP Survey of Refugees, a longitudinal survey of refugees, asylum seekers, and their household members in Germany (Brücker et al., 2016). The survey is conducted jointly by the Institute for Employment Research (IAB), the Research Centre of the Federal Office for Migration and Refugees (BAMF-FZ), and the Socio-Economic Panel (SOEP) at DIW Berlin. Its sampling frame is the CRF, also used in the main analysis. The survey is restricted to nationalities classified as having a good prospect of approval at the time of sampling, and oversampling was applied to women and persons aged over 30.

For analysis of mechanisms I use the German Socio-Economic Panel (SOEP) to investigate changes in beliefs at the site of arson attacks. The SOEP is a nationally representative longitudinal household survey that has interviewed approximately 15,000 households and 30,000 individuals annually since 1984 (Goebel et al., 2019). The survey contains detailed information on demographics, socioeconomic status, and migration background, as well as attitudes toward immigration and perceived threats. Geocoded identifiers at the county level allow me to match respondents to immigration office counties. I use SOEP v41 (survey years 2015–2019) for the public attitudes analysis in Section 7.3.

Office-level aggregate statistics. I complement the individual-level CRF data with office-level aggregate statistics on first-instance asylum decisions, originally compiled by Barreto et al. (2022). These data record the total number of decisions and their outcomes for each immigration office branch, nationality, and year cell, covering the period 2010–2020. Because the statistics tabulate first-instance decisions only and exclude appeals, they allow a direct test of whether arson attacks altered the content of initial decisions rather than the downstream appeal process. The aggregation, however, means that applicant-level covariates are unavailable and applicant behavior cannot be separated from officer behavior, so I use these data as a robustness check rather than as a primary specification (Section 6.1.1).

News data. To measure the media salience of arson attacks, I use raw LexisNexis article data originally collected by Graeber et al. (2026) to retrieve full-text news articles from major German newspapers, applying keyword filters for refugee shelter and arson terms. Articles are matched to attacks using a spatial window of 25 km and a temporal window of ± 7 days from the attack date, with manually verified overrides for attacks where geographic identifiers in the press coverage are imprecise. The resulting article count serves as the salience measure in the survey-based analysis of news amplification (Section 7.2).

5 Empirical Strategy

5.1 Identification Strategy

The CRF records only the final legally binding asylum status, observed after any appeals have concluded. Approved applications resolve at the time of the first-instance decision, but rejected applications may be appealed and appear in the data only much later. Using the date of the final outcome to define pre- and post-treatment periods would therefore conflate the *content* of decisions with their *timing*: a post-attack decline in observed positive outcomes could reflect either fewer approvals or a shift in when cases resolve.

To separate these two channels, I define the outcome as the monthly approval indicator

$$Y_{it} = \mathbf{1}(\text{applicant } i \text{ receives a positive decision in month } t) \quad (1)$$

with $Y_{it} = 0$ for all months if no approval is recorded within 12 months of application, regardless of whether the applicant was rejected, is still pending, or has entered the appeals process. This creates a person $\times$ month panel in which each applicant contributes observations from the month of application until either a positive decision or the 12-month censoring point.

Because approvals are not appealed, any first-instance approval is observed at its true decision date, even though rejections are not directly observed at the correct time. Among

still-pending applicants, the expected approval indicator decomposes as:

$$\mathbb{E}[Y_{it}] = \mathbb{E}\left[\underbrace{\lambda_{it}}_{\text{decision rendering}} \cdot \underbrace{Z_i}_{\text{approval outcome}}\right] \quad (2)$$

where λ_{it} is the probability that a first-instance decision of any kind is rendered for applicant i in month t among still-pending applicants, and $Z_i \in \{0, 1\}$ indicates whether the decision is positive. The treatment effect of interest is the change in $\mathbb{E}[Z_i]$. Recovering this from the observed left-hand side requires assumptions that discipline λ_{it} and the composition of the at-risk population. The design rests on three assumptions.

Assumption ID.1 (Exogenous applicant composition). Conditional on place of residence, the allocation of asylum applicants across processing branch offices is orthogonal to treatment status. For applications submitted prior to treatment, office assignment is predetermined and plausibly uncorrelated with treatment. For applications submitted after treatment, I assume that the flow of new filings to each office does not respond to nearby arson attacks. This is plausible because applicants are assigned to offices administratively and cannot choose their processing location. Figure B.1 tests this assumption by showing that predicted approval rates, constructed from pre-treatment cell means, do not shift around treatment timing.

Assumption ID.2 (Appeals fall outside the observation window). The combined first-instance processing time and appeals duration exceeds the 12-month observation window, so the only positive outcomes observed within that window are first-instance approvals: no successful appeal contaminates Y_{it} . This is empirically supported by institutional processing times: average asylum court proceedings took 12.5 months in 2018, on top of several months of first-instance immigration office processing.⁴

Assumption ID.3 (Content, not timing). Exposure to a nearby arson attack shifts the expected approval outcome $\mathbb{E}[Z_i]$ but not the decision-rendering hazard λ_{it} . Let D indicate whether a branch office has been exposed to a nearby arson attack. ID.3 requires that, within calendar month $\times$ duration $\times$ nationality cells:

$$\mathbb{E}[\lambda \mid D = 1] = \mathbb{E}[\lambda \mid D = 0] \quad (3)$$

so that λ_{it} is mean-independent of treatment status and any difference in $\mathbb{E}[Y_{it}]$ across treatment arms isolates the shift in $\mathbb{E}[Z_i]$, not in the timing of decisions. I test this assumption directly using office-level aggregate statistics in Section 6.1.1 and find no evidence that the decision-rendering hazard changed at treated offices.

⁴Source: parliamentary interpellations on asylum processing durations (Schwerpunktfragen zur Asylverfahrensdauer, BT-Drs. 20/5709).

Estimand. Under Assumptions ID.1–ID.3, the difference-in-differences coefficient identifies:

$$\beta = \mathbb{E}[\lambda] \cdot (\mathbb{E}[Z \mid D = 1] - \mathbb{E}[Z \mid D = 0]) \quad (4)$$

the shift in the approval content of decisions at exposed offices, scaled by the average decision-rendering hazard $\mathbb{E}[\lambda]$. The factorization from Equation 2 to 4 requires that λ_{it} and Z_i are conditionally independent given treatment status and cell membership: ID.3 ensures λ_{it} does not vary with D , and ID.1 ensures the composition of Z_i does not vary with D within cells, so the expectation of their product separates. β therefore isolates the change in approval rates for treated applications.

Treatment coding. A branch office enters the post-first-attack regime upon the first arson attack occurring within a 25 km radius. Subsequent attacks during the sample period are not coded separately, so the post-period coefficients reflect the joint effect of the first attack and any subsequent nearby incidents at the same office. This simplifies the event-study to a single cutoff per office but means the design does not, by itself, decompose the initial shock from cumulative exposure.

From office-level estimate to officer discretion. The office-level estimate could reflect three channels beyond individual caseworker discretion: differential managerial directives at treated offices, differential staffing or case-assignment responses, or applicant-side confounders operating within the office’s local ecosystem. Because asylum law is federal and immigration office headquarters issues centralized operational guidelines that apply identically to all 48 branch offices, responses that differentially affect offices such as increased oversight would have to be made explicit. The direct channel through which local conditions can enter adjudication is the individual officer’s exercise of discretion, in assessing credibility, weighing evidence, and determining protection outcomes. Appendix B.1 provides an extended discussion of institutional responses and other threats.

Threats to identification. Three first-order concerns bear on the design. First, arson attacks cluster in areas with rising anti-immigrant sentiment, raising the possibility that offices are already responding to the same underlying sentiment before the attack occurs. The fixed effects absorb nationwide sentiment shifts, time-invariant local attitudes, and level differences across applicants’ residence counties; a robustness specification additionally absorbs county-level time-varying shocks (Section 6.1). What remains unabsorbed is time-varying sentiment specific to the branch office’s locality. The event study examines whether pre-treatment coefficients reveal differential trends at treated offices; the four pre-treatment coefficients are individually and jointly insignificant, though with a five-month pre-window, the test has limited power to detect violations that would

meaningfully bias the post-treatment estimate (Roth, 2022). The design also requires a no-anticipation assumption: officers do not adjust decisions in response to attacks that have not yet occurred. Given that arson attacks are criminal acts without public advance notice, anticipation effects are unlikely.

Second, attacks may affect decisions at distant offices through news coverage or internal communication, violating the stable unit treatment value assumption. Two types of spillover are relevant: geographic spillover via news coverage affecting applicants' locations, and office-to-office spillover via internal immigration office communication. I probe the former by restricting the sample to applicants whose residence county differs from their processing office county, and the latter using alternative control groups. Section 6.1 presents these tests, and Appendix B.1 provides extended discussion of each threat, including selective attrition, institutional responses, and legal representation.

Third, the immigration office's large-scale hiring during 2015–2017 raises the possibility of differential staffing responses at treated offices, which would conflate behavioral change among incumbent officers with compositional effects of new hires. Without officer-level identifiers in the CRF, I cannot directly separate these channels. Appendix B.1 addresses this concern using register attrition data (Table B.1), finding no evidence of differential attrition across all counties or the subset hosting immigration offices. Section 7.1 provides complementary evidence: the arson effect is concentrated among cases adjudicated through personal interviews, while applicants processed via written procedures at the same office show no comparable response, consistent with a within-officer behavioral channel rather than a purely compositional one.

5.2 Estimation Strategy

I estimate two complementary specifications on a person $\times$ month panel: a collapsed difference-in-differences and a dynamic event study. Each is implemented using both a heterogeneity-robust stacking estimator (preferred) following Baker et al. (2022) and a conventional TWFE specification (benchmark). The outcome Y_{irt} is a binary indicator equal to one if applicant i , residing in county r and processed by branch office d , receives a positive first-instance decision in calendar month t , and zero otherwise. I compare treated offices against not-yet-treated and never-treated offices. Treatment timing varies across the 27 treated offices, creating a staggered adoption design.⁵

Collapsed difference-in-differences. For each treatment cohort g , defined by the calendar month D_g of the first arson attack within 25 km, I construct a sub-experiment pairing the offices treated at D_g with all not-yet-treated and never-treated offices, restricted to the event window $[D_g - 5, D_g + 11]$. Stacking the sub-experiments yields:

⁵Similar identification strategies have been applied to assess mortality rates instead of approval rates (Miller et al., 2021).

$$Y_{irdts} = \beta \cdot \text{Treated}_d \cdot \mathbf{1}(t \geq D_{g(s)}) + \alpha_{d,s} + \mu_{r,s} + \delta_{t,\tau,n,s} + X_i' \Gamma + \varepsilon_{irdts} \quad (5)$$

where Treated_d equals one for offices in treatment cohort $g(s)$ and zero for not-yet-treated and never-treated offices, s indexes sub-experiments, τ denotes months since application (duration), n denotes nationality, $g(s)$ maps each sub-experiment to its treatment cohort's attack month, and Γ is the coefficient vector on individual controls X_i . All fixed effects are interacted with the sub-experiment indicator: $\alpha_{d,s}$ (branch office $\times$ sub-experiment), $\mu_{r,s}$ (residence county $\times$ sub-experiment), and $\delta_{t,\tau,n,s}$ (calendar month $\times$ duration $\times$ nationality $\times$ sub-experiment), ensuring that each treated cohort is compared only against its own clean control group. X_i includes gender, family status, birth year, and minor status. Standard errors are clustered at the branch office level. The treatment coefficient β estimates the average change in the monthly probability of a positive decision across the twelve post-treatment months at treated offices. The 43 offices provide the identifying variation, with 27 experiencing a first attack within 25 km during the sample period.

Event-study specification. To trace the temporal profile of the treatment effect and check pre-trends, I replace the post-treatment indicator with event-time dummies:

$$Y_{irdts} = \sum_{\substack{k=-5 \\ k \neq -1}}^{11} \beta_k \cdot \text{Treated}_d \cdot \mathbf{1}(t - D_{g(s)} = k) + \alpha_{d,s} + \mu_{r,s} + \delta_{t,\tau,n,s} + X_i' \Gamma + \varepsilon_{irdts} \quad (6)$$

The coefficients β_k identify differential approval probabilities, within applications that share the same branch office, residence county, and calendar month $\times$ time since application $\times$ nationality cell, between applicants processed by offices exposed to a nearby arson attack k months prior and those processed by not-yet- or never-treated offices in the same sub-experiment.

6 Results

The preferred stacking estimator yields a reduction of 0.6 percentage points in the monthly probability of receiving any form of protection at offices exposed to a nearby arson attack (Table 2 column 1), corresponding to a 23% decline relative to the sample mean of 2.6%. Using the not-yet-treated control group, the estimate is -1.4 percentage points and significant at the 10 percent level, consistent in sign but noisier due to the smaller comparison sample.




Figure 3: Effect of Arson Attacks on the Monthly Probability of a Positive Asylum Decision

Notes: Stacking estimator (Equation 6) following Baker et al. (2022). Coefficients show change in the probability of receiving any form of protection relative to the month before treatment ($k = -1$). Vertical bars indicate 95% confidence intervals. Standard errors clustered at the branch office level. Sample: 2015–2018. Source: AZR Research Dataset 2023; parliamentary interpellations (arson events).

Figure 3 presents the dynamic event-study coefficients from the stacking estimator. The pre-treatment coefficients are statistically insignificant, supporting the parallel-trends assumption. Effects become statistically significant at $k = +2$ and increase steadily, peaking at months 7–8 before fading after month 9. This delayed onset is consistent with the institutional lag between asylum interviews, where adjudication is concentrated, and the recording of decisions (see Appendix Figures A.2 and A.3 for processing durations by stage). The interview of applicants is the moment at which officer discretion is highest and where treatment is most likely to affect it. Arson attacks that precede the interview are most likely to affect the outcome of an application, explaining the delayed onset in approval probabilities. The period in which the effects are highest (months 5–8, peaking at 7–8) mirrors the average delay between interview and the rendering of a decision (see Figure A.3 in the Appendix).

Table 2: Effect of arson attacks on the monthly probability of a positive asylum decision (collapsed post-period)

	(1) Never-treated	(2) Not-yet-treated
Post $\times$ Treated	-0.006** (0.003)	-0.014* (0.008)
Mean dep. var.	0.026	0.016
Observations	2,682,401	857,847
Persons	95,320	61,722
Office $\times$ sub-exp. FE	Yes	Yes
Res. county $\times$ sub-exp. FE	Yes	Yes
Month $\times$ duration $\times$ nat. $\times$ sub-exp. FE	Yes	Yes

Notes: Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Column (1): never-treated comparison group; the point estimate corresponds to a 23% decline relative to the sample mean. Column (2): not-yet-treated comparison group. Standard errors clustered at the branch office level in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

6.1 Robustness

The most direct robustness check restricts the sample to applicants whose residence county differs from their processing office county; the effect persists, reinforcing the interpretation of an office-side channel rather than an applicant-side response to local violence.⁶ The estimate is also stable across five alternative fixed-effect structures (Appendix Table C.2), ranging from a parsimonious specification that enters duration, calendar month, and nationality effects separately to richer specifications that interact residence county with calendar month or add office-state $\times$ calendar month effects. All five specifications yield significant negative coefficients of similar magnitude.

Assumption ID.2 requires that no successful appeal is observed within the observation window. Shortening the maximum observation window from 12 to 9, 6, or 3 months after application tightens this assumption, as shorter windows make it increasingly unlikely that any court decision contaminates the outcome. All three restricted windows yield statistically significant estimates that are monotonically larger than the baseline (Appendix Table C.3), confirming that the results are not driven by appeals entering the observation window and that the effect is concentrated in early decision periods.

6.1.1 Office-Level Aggregate Statistics

A complementary check on the main results uses the office-level aggregate statistics described in Section 4.3.⁷ Because the data are aggregated at the office level, they lack the

⁶On the different-county subsample (79,137 persons), the estimate is -0.7 percentage points (SE 0.3), significant at the 1% level.

⁷I thank Martin Lange and co-authors for sharing the office-level aggregate data from Barreto et al. (2022).

demographic and processing-stage detail of the CRF sample used in the main analysis. Applicant behavior cannot be disentangled from immigration officer behavior at this level of aggregation. I therefore use these data as a robustness check rather than a primary specification.

The office-level data are well suited to verify two assumptions underlying the main identification. The aggregate counts include only first-instance decisions and exclude appeals, so any effect detected here cannot be driven by differential appeal behavior across treated and untreated offices. The data also allow a direct test of Assumption ID.3. If arson attacks shifted the timing of decisions rather than their content, the number of decisions at treated offices should decline.

I apply a stacking estimator to the office-level panel, using a continuous exposure measure that exploits within-year timing variation. For each treatment cohort, I pair treated offices with never-treated offices over pre-treatment and treatment years, interacting all fixed effects with the stack indicator. The exposure variable $E_d = (13 - m_d)/12$, where m_d is the calendar month of office d 's first attack, assigns higher exposure to offices attacked earlier in the year (a January attack yields $E_d = 1$; a December attack yields $E_d \approx 0.08$). The sample covers 2015 to 2018, and excludes Syrian, Iraqi, and Eritrean nationalities, which received almost universal approval in 2015 through written procedures and therefore lack a meaningful pre-treatment period.

Table 3 reports estimates from the stacking specification

$$Y_{dnys} = \alpha_{d,s} + \mu_{n,y,s} + \beta E_{ds} + \varepsilon_{dnys}, \quad (7)$$

where Y_{dnys} is the outcome for office d , nationality n , year y , and stack s . $\alpha_{d,s}$ are office $\times$ stack fixed effects and $\mu_{n,y,s}$ are nationality $\times$ year $\times$ stack fixed effects. E_{ds} equals the exposure measure for treated offices in their treatment year and zero otherwise. Columns (1) and (2) use the protection rate as the outcome, weighted by the number of decisions. Columns (3) and (4) use log decisions to test whether attacks shifted the volume of decisions.

The exposure coefficient on the protection rate is negative in both specifications but statistically significant only after adding state $\times$ year $\times$ stack fixed effects: the estimate in the parsimonious specification is -0.036 (SE 0.024, column 1), while adding state $\times$ year $\times$ stack fixed effects yields -0.068 ($p < 0.01$, column 2). The sensitivity to these additional controls likely reflects the coarseness of the yearly panel: with only 52 offices observed over four years, uncontrolled state-level trends in approval rates—driven, for example, by differential caseload composition across states—can attenuate or mask the exposure effect. The monthly person-level CRF data used in the main analysis absorb these trends through richer fixed-effect structures and yield a precisely estimated effect without relying on state $\times$ year controls. Despite the limited precision, the direction and magnitude of the office-level estimates are consistent with the main finding, and the estimate is obtained from an

independent dataset that records only first-instance decisions and excludes appeals. The log-decisions placebo in columns (3) and (4) shows no significant effect on the volume of decisions at treated offices, consistent with Assumption ID.3: arson attacks shifted the content of decisions, not their timing or volume.

Table 3: Office-Level Aggregate Statistics

	(1)	(2)	(3)	(4)
	Protection rate		Log decisions	
Exposure	-0.036 (0.024)	-0.068*** (0.026)	0.205 (0.222)	0.085 (0.353)
Mean dep. var.	0.177	0.177	3.756	3.756
Observations	2,745	2,745	2,745	2,745
Offices	52	52	52	52
Treated	26	26	26	26
State $\times$ year $\times$ stack FE	No	Yes	No	Yes

Notes: Stacking estimator following [Baker et al. \(2022\)](#) with continuous exposure $E_d = (13 - \text{attack month})/12$. Cols. (1)–(2): protection rate (weighted by decisions); Cols. (3)–(4): log decisions (unweighted). Sample: 2015–2018, excluding Syrian, Iraqi, and Eritrean nationalities. Standard errors are clustered at office level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

7 Discussion

7.1 Personal Interview as the Site of Discretion

The main results rely on the institutional argument that officer discretion is concentrated at the personal interview. A separate policy change provides direct evidence for this claim. During 2015, the immigration office processed claims from Syrian, Eritrean, and Iraqi-minority applicants through simplified written procedures without individual interviews. Beginning in 2016, personal interviews were reinstated: applicants who arrived after 1 January 2016 were required to undergo a personal interview, and immigration office headquarters later extended this requirement to all persons filing from 17 March 2016, regardless of entry date.⁸ This creates within-office variation in interview format among applicants processed at the same branch office during the same period.

The specification, estimated on the three major refugee-sending nationalities (Syria, Afghanistan, Iraq), absorbs branch-office fixed effects (interacted with the sub-experiment indicator) and allows county-level trends and the baseline hazard to differ across interview formats through personal-interview-by-county-by-calendar-month and personal-interview-by-duration-by-calendar-month-by-nationality fixed effects. The post-attack coefficient

⁸The personal interview indicator equals one if the applicant arrived after 1 January 2016 or filed after 17 March 2016. Iraqi applicants, who were only partially subject to the original written procedure, are included in the analysis.

is identified from within-office variation over time and captures the arson effect on applicants processed via written procedure. The interaction captures the additional effect for applicants adjudicated through a personal interview. The reform that reinstated personal interviews bundled this change with other policy measures;⁹ the interview-format-by-county-by-calendar-month fixed effects absorb any uniform effects of the broader policy package, isolating the interview-format margin.

Table 4 interacts the post-arson-attack indicator with the personal interview indicator, testing directly whether the arson effect operates through the interview stage. Column (1) uses any form of protection as the outcome; column (2) uses full refugee status. The post-attack coefficient, which captures the arson effect on applicants processed via written procedure, is close to zero and statistically insignificant in both columns: non-interview cases show no response to nearby attacks. The interaction term is large, negative, and highly significant, indicating that the arson effect is concentrated entirely among applicants whose claims were adjudicated through a personal interview. This confirms that the interview is the channel through which local hostility enters adjudication: officers who never meet applicants face-to-face do not adjust their decisions in response to arson exposure, while those conducting personal interviews do.

Table 4: Arson Effect by Interview Format

	(1) Any protection	(2) Full refugee
Post-Attack	0.004 (0.014)	0.000 (0.013)
Post-Attack $\times$ Personal Interview	-0.047*** (0.010)	-0.054*** (0.010)
Mean dep. var.	0.105	0.074
Observations	1,872,610	2,077,281
Persons	71,309	71,309

Notes: Post-arson-attack indicator and its interaction with a personal interview indicator. Column (1): dependent variable is any form of protection; column (2): full refugee status. Sample restricted to the three major refugee-sending nationalities (Syria, Afghanistan, Iraq) with positive baseline acceptance rates, applications filed January 2015–December 2016. All specifications include branch office $\times$ sub-experiment, personal interview $\times$ county $\times$ calendar month $\times$ sub-experiment, and personal interview $\times$ duration $\times$ calendar month $\times$ nationality $\times$ sub-experiment fixed effects, as well as demographic controls (gender, family status, birth year, minor status). Standard errors clustered at the branch office level in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

⁹The *Asylpaket II* (February 2016) combined the reinstatement of personal interviews with safe-country designations and family reunification restrictions.

7.2 Survey evidence on approval rates

The preceding results rely on administrative data that does not include information on interview dates. An independent test uses individual-level data from the IAB-BAMF-SOEP refugee survey, where respondents report when their interview took place and whether their claim was approved. The main shortcoming of the survey data is that it does not include information on the assigned immigration office. The treatment definition therefore has to be based on the residence of an applicant. I exclude any individuals with an arson attack within their county (of which there are 423) and assign their applications as treated if there is an arson attack within the coarser planning region (96 regions). The survey is not fully representative of all asylum applications as it is restricted to nationalities classified as having a good prospect of approval at the time of sampling. On top of the nationality dimension, oversampling was also applied to women and persons aged over 30. The cross-sectional specification estimates:

$$Y_i = \beta_1 D_{r(i),t} + \beta_2 D_{r(i),t} \times N_{r(i),t} + X_i' \Gamma + \mu_t + \varepsilon_i \quad (8)$$

where Y_i equals one if respondent i reports a positive asylum decision. $D_{r(i),t}$ indicates whether any arson attack occurred in the respondent's planning region r (excluding the respondent's own county) during month t , defined alternately by the interview month or the decision month. $N_{r(i),t}$ is the Nexis newspaper article count for attacks in the region, standardized to one standard deviation. X_i includes personal characteristics, and μ_t absorb monthly fixed effects. Standard errors are clustered at the planning region level. I also estimate an event-study variant replacing $D_{r(i),t}$ with relative-month indicators $\mathbf{1}(t - D_{r(i)} = k)$ for $k \in \{-5, \dots, 5\}$.

Table 5 tests whether the effect operates through interview-month or decision-month exposure and whether news coverage amplifies it. Columns (1)–(2) define treatment by the interview month; columns (3)–(4) by the decision month. Binary arson exposure at the interview month significantly reduces approval probability (column 1). Adding an interaction with newspaper coverage (column 2), the interaction is negative and significant, indicating that attacks receiving more press coverage produce larger effects, while the binary coefficient attenuates. When treatment is instead defined by the decision month (columns 3–4), the binary coefficient is close to zero and insignificant, consistent with the personal interview—where personal discretion of immigration officers is highest—as the site where local conditions enter adjudication.

Figure 4 presents event-study coefficients from this specification. Panel (a) plots the binary arson effect. Pre-treatment coefficients are close to zero and statistically insignificant. The approval probability drops sharply in the month of the arson attack and remains negative, though imprecisely estimated, at longer horizons. Panel (b) plots the interaction between binary arson exposure and the news article count (per one standard deviation). Pre-treatment interaction coefficients are close to zero. After treatment, the

Table 5: Arson Exposure and Approval Probability (SOEP)

	(1)	(2)	(3)	(4)
Arson	-0.049** (0.020)	-0.029 (0.024)	0.011 (0.018)	0.008 (0.021)
Arson $\times$ News		-0.010* (0.005)		0.002 (0.002)
Observations	3,044	3,044	2,917	2,917
Timing	Interview	Interview	Decision	Decision
Mean dep. var.	0.859	0.859	0.863	0.863

Notes: Linear probability models. Dependent variable: individual asylum approval (SOEP self-report). “Arson”: binary indicator for any arson attack in the respondent’s planning region (excluding own county) during the relevant month. “Arson $\times$ News”: interaction with Nexis newspaper article count (standardized to 1 SD). Columns (1)–(2) define treatment by interview month; columns (3)–(4) by decision month. All specifications include calendar-month fixed effects, controlling for gender, age at interview, and immigration office administrative approval rate. Standard errors clustered at the planning-region (Raumordnungsregion) level in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

interaction turns negative, indicating that within the set of exposed respondents, those whose local attacks received more press coverage experienced larger declines in approval probability. The pattern is consistent with a transient salience channel in which media coverage amplifies the initial signal of local hostility.




Figure 4: Event study: arson exposure and individual approval (SOEP)

Notes: Panel (a): coefficients on relative-month indicators from an individual-level regression of asylum approval. Treatment: nearest arson attack in the same ROR, excluding the respondent’s own county. Reference category: applicants with no arson within ± 5 months (Emeriau, 2024). Panel (b): interaction coefficients from separate cross-sectional regressions at each offset k , interacting binary arson exposure with the Nexis article count (standardized to 1 SD). Controls: gender, age at interview, immigration office administrative approval rate. Calendar-month fixed effects. Standard errors clustered at the planning region level. Error bars show 95% confidence intervals. Data: IAB-BAMF-SOEP refugee survey, 2016–2019; arson events from parliamentary interpellations; Nexis newspaper coverage.

These results triangulate the main administrative findings using an independently collected survey. The SOEP sample is small ($N = 3,044$ under interview-month timing) and treatment is defined at a different geographic level (planning region rather than branch office), so the estimates are noisier than the main specification. The direction of the approval decline replicates across data sources, while the media-salience channel is identified only in the survey and should be read as suggestive: the interaction rests on a single cross-sectional specification without region fixed effects, comparing across planning regions with and without a contemporaneous attack rather than tracking within-region changes.

7.3 Beliefs

I use survey data from the SOEP to test how attacks against refugees shifted beliefs near immigration offices. Only counties that host an immigration office branch enter the sample. I construct a stacked difference-in-differences design paralleling the main specification, comparing SOEP respondents in treated immigration office counties against respondents in never-treated immigration office counties. The event-study specification is:

$$\tilde{W}_{ics} = \sum_{\substack{k=-3 \\ k \neq -1}}^8 \beta_k \cdot \mathbf{1}(k_s = k) \cdot D_c + \phi \tilde{W}_{i,s-1} + \alpha_{g,s} + \mu_{k,s} + \varepsilon_{ics} \quad (9)$$

where $\tilde{W}_{ics}$ is the worry outcome for respondent i in county c and sub-experiment s , standardized by the pre-treatment standard deviation. D_c indicates whether the respondent resides in a county whose immigration office was exposed to a nearby arson attack; all respondents in the sample reside in counties that host an immigration office branch. The coefficients β_k are identified by comparing the evolution of worry outcomes between treated and never-treated office counties around each attack. Because interview dates vary across respondents within each SOEP wave, the event window typically contains only one observation per respondent, so the identifying variation is cross-sectional—comparing respondents interviewed before the attack to those interviewed after, in treated versus control counties. $\tilde{W}_{i,s-1}$ is the lagged outcome from the previous survey wave, which exploits the panel dimension of the SOEP to control for prior attitudes. The fixed effects $\alpha_{g,s}$ (treatment group $\times$ sub-experiment) and $\mu_{k,s}$ (month bin $\times$ sub-experiment) absorb group-level means and common time patterns within each sub-experiment. Standard errors are clustered at the county level.

Figure 5 presents event-study coefficients for respondents residing in treated immigration office counties, expressed in pre-treatment standard deviations. Pre-treatment coefficients are statistically insignificant for both outcomes. If anything, immigration worry displays a declining trend before treatment, making the post-attack increase a conservative estimate. The two outcomes display distinct temporal patterns. Xenophobia worry spikes sharply in the month of the attack but drops back to near zero by month 1 and remains flat

thereafter. Immigration worry rises more modestly at impact but stays elevated and continues to grow over subsequent months, though confidence intervals widen substantially at longer horizons. Arson attacks trigger an immediate concern about xenophobic hostility that fades as the event recedes from public attention, while worries about immigration accumulate as the local salience of the refugee issue persists.

The short-lived increase in concern for refugee safety aligns with Freitas-Monteiro and Prömel (2024), who show a short-term increase in pro-refugee beliefs after large far-right protests at the national level, and with Emeriau (2024), who documents a short-run empathy response in approval rates after migrant shipwrecks. The more persistent pattern of immigration worry is also consistent with the voting responses to far-right terrorist attacks documented by Sabet et al. (2023) and with the longer-lasting decrease in approval rates found in the main results.




Figure 5: Event study: public attitudes in treated counties

Notes: Stacked difference-in-differences event-study coefficients. Dependent variables are SOEP worry items (1–3 scale, increasing in concern), standardized by the pre-treatment standard deviation. Treatment: first arson attack in the respondent’s county of residence. Control group: respondents in never-treated immigration office counties. Reference period: month -1 . All specifications include cohort $\times$ treated and cohort $\times$ month-bin fixed effects, controlling for lagged immigration worry, xenophobia worry, and social cohesion worry. Standard errors clustered at the county level. Error bars show 95% confidence intervals. Data: SOEP v41, 2015–2019.

8 Conclusion

This paper uses geographic separation between processing offices and applicant residences in the German asylum system to test whether far-right violence compromises high-stakes administrative decisions. Arson attacks against refugee accommodations reduce the monthly approval probability at nearby branch offices by 0.6 percentage points (23 percent relative to the mean), with larger effects for attacks involving risk to human life and in regions where far-right support is not already entrenched. A within-office test confirms the personal interview as the site of discretion: caseworkers who meet applicants face-to-face adjust their decisions after nearby attacks, while those deciding cases through written procedures at the same office do not. Far-right violence can, in sum, shift the outcomes of a process designed to reflect the merits of each individual claim.

What makes these findings distinct is that the hostility signal is *in-group* violence against a vulnerable out-group. Unlike out-group-perpetrated terrorism (Shayo and Zussman, 2011; Brodeur and Wright, 2019), arson attacks on refugee housing offer no rational basis for updating threat assessments, yet the bureaucratic response mirrors them. Survey evidence indicates that compliance dominates resistance as responses to political violence. The hostile signal accumulates in the local environment while sympathy for victims fades, providing a plausible attitudinal pathway through which the political climate enters the adjudication room. The results point toward concrete safeguards. Structured interviews with standardized question catalogs could limit the scope for subjective assessment. Systematic training and feedback of the kind Emeriau (2023) shows can reduce bias among asylum officers would address the individual-level channel, while rotation of caseworkers across regions could insulate offices from persistent local conditions. More broadly, wherever public servants exercise discretion within the boundaries of formal rules (Lipsky, 2010), their decisions may absorb the political climate of their environment—a vulnerability that extends well beyond immigration policy.

References

- Aidt, T. S. and R. Franck (2015). Democratization Under the Threat of Revolution: Evidence From the Great Reform Act of 1832. *Econometrica* 83(2), 505–547. 1
- Aksoy, C. G., P. Poutvaara, and F. Schikora (2023, September). First time around: Local conditions and multi-dimensional integration of refugees. *Journal of Urban Economics* 137, 103588. 2.1
- Arnold, D., W. Dobbie, and C. S. Yang (2018, November). Racial Bias in Bail Decisions. *The Quarterly Journal of Economics* 133(4), 1885–1932. 1
- Baker, A. C., D. F. Larcker, and C. C. Y. Wang (2022, May). How much should we trust staggered difference-in-differences estimates? *Journal of Financial Economics* 144(2), 370–395. 5.2, 3, 3
- BAMF-Forschungsdatenzentrum (2023). AZR Research Dataset 2021 (Version 1.1, Dataset by BAMF-FDZ). Nuremberg. Federal Office for Migration and Refugees. <https://doi.org/10.48570/bamf.fdz.azr.on.2021.1.1>. 12
- Barreto, C., P. Berbée, K. Gallegos Torres, M. Lange, and K. Sommerfeld (2022, April). The Civic Engagement and Social Integration of Refugees in Germany. *Nonprofit Policy Forum* 13(2), 161–174. 4.3, 7
- Bolhaar, J., N. Ketel, and B. Van Der Klaauw (2020, March). Caseworker’s discretion and the effectiveness of welfare-to-work programs. *Journal of Public Economics* 183, 104080. 1
- Brodeur, A. and T. Wright (2019, December). Terrorism, immigration and asylum approval. *Journal of Economic Behavior & Organization* 168, 119–131. 1, 8
- Brücker, H., N. Rother, and J. Schupp (2016). *IAB-BAMF-SOEP-Befragung von Geflüchteten: Überblick und erste ergebnisse*, Volume 29. DEU. 4.3
- Bursztyn, L., G. Egorov, and S. Fiorin (2020, November). From Extreme to Mainstream: The Erosion of Social Norms. *American Economic Review* 110(11), 3522–3548. 1
- Chen, D. L. and M. Loecher (2025, June). Mood and the malleability of moral reasoning: the impact of irrelevant factors on judicial decisions. *Journal of Behavioral and Experimental Economics* 116, 102364. 1
- Chen, D. L., T. J. Moskowitz, and K. Shue (2016, August). Decision Making Under the Gambler’s Fallacy: Evidence from Asylum Judges, Loan Officers, and Baseball Umpires*. *The Quarterly Journal of Economics* 131(3), 1181–1242. 1

- Chenoweth, E. and M. J. Stephan (2011). *Why Civil Resistance Works: The Strategic Logic of Nonviolent Conflict*. Columbia University Press. Google-Books-ID: kGW8scops3GcC. 1, 3
- Danziger, S., J. Levav, and L. Avnaim-Pesso (2011, April). Extraneous factors in judicial decisions. *Proceedings of the National Academy of Sciences* 108(17), 6889–6892. 1
- Dobbie, W., J. Goldin, and C. S. Yang (2018, February). The Effects of Pre-Trial Detention on Conviction, Future Crime, and Employment: Evidence from Randomly Assigned Judges. *American Economic Review* 108(2), 201–240. 1
- Emeriau, M. (2023, October). Learning to be Unbiased: Evidence from the French Asylum Office. *American Journal of Political Science* 67(4), 1117–1133. 1, 3, 8
- Emeriau, M. (2024, October). Victim or Threat? Shipwrecks, Terrorist Attacks, and Asylum Decisions in France. *American Journal of Political Science* 68(4), 1187–1204. 1, 2.2, 3, 4, 7.3
- Eren, O. and N. Mocan (2018, July). Emotional Judges and Unlucky Juveniles. *American Economic Journal: Applied Economics* 10(3), 171–205. 1
- Feinberg, M., R. Willer, and C. Kovacheff (2020). The activist’s dilemma: Extreme protest actions reduce popular support for social movements. *Journal of Personality and Social Psychology* 119(5), 1086–1111. 1, 3
- Freitas-Monteiro, T. and C. Prömel (2024, December). Local far-right demonstrations and nationwide public attitudes towards migration. *Journal of Economic Behavior & Organization* 228, 106766. 1, 3, 7.3
- Gassebner, M., R. Jong-A-Pin, and J. O. Mierau (2008). Terrorism and electoral accountability: One strike, you’re out! *Economics Letters* 100(1), 126–129. 1
- Gleiser, P. and T. Hinz (2024). Datenreport zum AZR-Forschungsdatensatz 2023. Technical Report Veröffentlichungen des BAMF-FDZ 01/2024, Bundesamt für Migration und Flüchtlinge, Nürnberg. 4.2
- Gleiser, P. and C. Salas Escobar (2024). AZR Research Dataset 2023 (Version 1.0, Dataset by BAMF-FDZ). Nuremberg. Federal Office for Migration and Refugees. <https://doi.org/10.48570/bamf.fdz.azr2023.on.1.0>. 4.2, 12
- Goebel, J., M. M. Grabka, S. Liebig, M. Kroh, D. Richter, C. Schröder, and J. Schupp (2019, April). The German Socio-Economic Panel (SOEP). *Jahrbücher für Nationalökonomie und Statistik* 239(2), 345–360. 4.3
- Gould, E. D. and E. F. Klor (2010). Does Terrorism Work? *The Quarterly Journal of Economics* 125(4), 1459–1510. 1

- Gould, E. D. and E. F. Klor (2016, November). The Long-run Effect of 9/11: Terrorism, Backlash, and the Assimilation of Muslim Immigrants in the West. *The Economic Journal* 126(597), 2064–2114. 1
- Graeber, D., L. Meister, and N. Murray (2026). Risk aversion in the shadow of terrorism. DIW Berlin. 4.3
- Grosjean, P., F. Masera, and H. Yousaf (2023, February). Inflammatory Political Campaigns and Racial Bias in Policing*. *The Quarterly Journal of Economics* 138(1), 413–463. 1, 3
- Gundacker, L., Y. Kosyakova, and G. Schneider (2024, December). How regional attitudes towards immigration shape the chance to obtain asylum: Evidence from Germany. *Migration Studies* 13(1), mnae002. 1
- Hemker, J. and A. Rink (2017, October). Multiple Dimensions of Bureaucratic Discrimination: Evidence from German Welfare Offices. *American Journal of Political Science* 61(4), 786–803. 1
- Heyes, A. and S. Saberian (2019, April). Temperature and Decisions: Evidence from 207,000 Court Cases. *American Economic Journal: Applied Economics* 11(2), 238–265. 1
- Jaschke, P., S. Sardoschau, and M. Tabellini. NBER WORKING PAPER SERIES. 2.1
- Jomni, H. and N. Zakharov (2026, January). Do terrorist attacks polarize politicians? Evidence from the European parliamentary speeches on migration. *Journal of Peace Research* 63(1), 86–97. 1
- Larsen, E. G., D. Cutts, and M. J. Goodwin (2020). Do terrorist attacks feed populist Eurosceptics? Evidence from two comparative quasi-experiments. *European Journal of Political Research* 59(1), 182–205. 1
- Lipsky, M. (2010). *Street-Level Bureaucracy, 30th Ann. Ed.: Dilemmas of the Individual in Public Service*. Russell Sage Foundation. 1, 3, 8
- Maestas, N., K. J. Mullen, and A. Strand (2013, August). Does Disability Insurance Receipt Discourage Work? Using Examiner Assignment to Estimate Causal Effects of SSDI Receipt. *American Economic Review* 103(5), 1797–1829. 1
- Miller, S., N. Johnson, and L. R. Wherry (2021, August). Medicaid and Mortality: New Evidence From Linked Survey and Administrative Data*. *The Quarterly Journal of Economics* 136(3), 1783–1829. 5
- Rees, D. I. and B. Smith (2022). Terror attacks and election outcomes in Europe, 1970–2017. *Economics Letters* 219, 110770. 1

- Roth, J. (2022). Pretest with caution: Event-study estimates after testing for parallel trends. *American Economic Review: Insights* 4(3), 305–322. [5.1](#)
- Sabet, N., M. Liebald, and G. Friebel (2023). Terrorism and Voting: The Rise of Right-Wing Populism in Germany. *SSRN Electronic Journal*. [1](#), [3](#), [7.3](#)
- Shayo, M. and A. Zussman (2011, August). Judicial Ingroup Bias in the Shadow of Terrorism *. *The Quarterly Journal of Economics* 126(3), 1447–1484. [1](#), [8](#)
- Spirig, J. (2023, January). When Issue Salience Affects Adjudication: Evidence from Swiss Asylum Appeal Decisions. *American Journal of Political Science* 67(1), 55–70. [1](#), [3](#)
- Wasow, O. (2020, August). Agenda Seeding: How 1960s Black Protests Moved Elites, Public Opinion and Voting. *American Political Science Review* 114(3), 638–659. [3](#)

Appendix Contents

A	Institutional Details and Data	32
A.1	Institutional Details	32
A.2	Descriptive Statistics	35
B	Identification	39
B.1	Identification Threats	39
B.2	Balance Tests	41
C	Additional Results	44
C.1	Heterogeneous Effects	44
C.2	Robustness	44

A Institutional Details and Data

A.1 Institutional Details

Internal Quality Assurance. The immigration office operates a dual-review principle under which a second reviewer checks each asylum decision before dispatch. During the 2015–2016 caseload surge, internal oversight was substantially weakened: the central quality-assurance unit reviewed approximately 1% of the 282,700 decisions rendered in 2015, and the dual-review requirement was made formally binding by central directive only in September 2017.¹⁰ A mandatory random 10% pre-dispatch quality check was introduced in May 2018 following the discovery of procedural irregularities at the Bremen branch office. These dates imply that for much of the sample period (2015–mid-2017), the institutional constraints on individual officer discretion were weaker than the standing procedural framework would suggest.

¹⁰Source: parliamentary inquiries BT-Drs. 19/4853 and testimony by the immigration office staff council before the Interior Committee in 2018.

Table A.1: Types of International Protection in Germany

Protection Status	Requirements	Legal Basis	Residence Duration	Family Reunification	Permanent Residency
Asylum (Art. 16a GG)	Political persecution by state actors; must enter directly from a non-safe country	German Basic Law (Grundgesetz)	3 years	Yes (easier, no waiting period)	After 3 years
Refugee Protection	Persecution due to race, religion, nationality, political opinion, or social group	Geneva Refugee Convention, §3 AsylG	3 years	Yes (easier, no waiting period)	After 3 years
Subsidiary Protection	No individual persecution but serious harm (e.g., death penalty, torture, war-related threats)	§4 AsylG	1 year (extendable to 3)	Restricted (2-year waiting period, proof of financial means required)	After 5 years
National Ban on Deportation	Risk of inhumane treatment or life-threatening situation in home country, but no refugee/subsidiary protection status	§60 Abs. 5/7 AufenthG	1 year (renewable)	Restricted (must prove financial stability)	After 7 years

Notes: Overview of protection categories under German asylum law as of 2015–2018 (the sample period). Legal bases: German Basic Law (Art. 16a GG), Asylum Act (AsylG), and Residence Act (AufenthG). Residence permit durations and family reunification rules reflect the legal framework during the study period. Source: Federal Office for Migration and Refugees procedural guidelines.

Table A.2: Classification of Arson Offenses (German Criminal Code)

Category	StGB §	German Title		Official Title	English	Key Statutory Elements
<i>Simple arson</i>						
Simple	§ 306	Brandstiftung		Arson		Sets fire to or destroys buildings, plants, warehouses, motor vehicles, forests, or agricultural facilities belonging to another (1–10 years)
Simple	§ 306d	Fahrlässige Brandstiftung		Negligent arson		Acts negligently in cases under § 306(1) or § 306a(1), or causes danger under § 306a(2) by negligence (≤5 years)
<i>Aggravated arson</i>						
Aggravated	§ 306a	Schwere Brandstiftung		Aggravated arson		Sets fire to premises serving to accommodate people, churches, or premises where people are usually present (≥1 year)
Aggravated	§ 306b	Besonders schwere Brandstiftung		Especially aggravated arson		Causes serious health damage or places another in danger of death; acts to facilitate another offence or prevents extinguishing (≥2 years; ≥5 years for danger of death)
Aggravated	§ 308	Herbeiführen einer Sprengstoffexplosion		Causing explosion		Causes an explosion, in particular using explosives, and thereby endangers life, limb, or property of significant value (≥1 year)

Notes: Classification of arson offenses under the German Criminal Code (*Strafgesetzbuch*, StGB). Section titles and statutory elements are from the official English translation provided by the Federal Ministry of Justice (gesetze-im-internet.de). Penalty ranges in parentheses. Simple arson covers § 306 and § 306d. Aggravated arson covers § 306a, § 306b, and § 308 — the categories that endanger human life, used as the severity indicator in the mechanism analysis.

A.2 Descriptive Statistics

Table A.3: Descriptive Statistics: Asylum Applications

	Person-months	Unique persons
<i>Panel A: Sample construction</i>		
Raw panel	6,426,097	162,537
After regular application filter	3,572,100	96,132
After universal-rejection filter	3,564,490	95,945
Estimation sample: at-risk for first pos. decision	3,223,576	95,945
<i>Panel B: Treatment assignment in estimation sample</i>		
Never-treated (comparison, all periods)	2,321,652	33,938
Not-yet-treated (comparison, pre-treatment)	84,338	20,536
Treated (post first arson within 25 km)	798,190	61,682
<i>Panel C: Branch offices</i>		
Ever treated within sample window	27	
Never treated within sample window	16	
Total branch offices in estimation sample	43	

Notes: Construction of the estimation sample from the AZR Research Dataset 2023. Unit of observation is the applicant-month. The regular application filter excludes airport transit applications and the Nuremberg headquarters office. The universal-rejection filter drops nationalities with near-certain rejection rates ($>98\%$), for which officer discretion is minimal. Rows in Panel B are not mutually exclusive: applicants at offices treated mid-sample contribute pre-treatment person-months to the not-yet-treated row and post-treatment person-months to the treated row, so counts sum to more than the estimation-sample total. Panel C reports branch offices in the estimation sample (43 of 48 total). Source: AZR Research Dataset 2023; immigration office branch-office register; parliamentary interpellations (arson events).

Table A.4: Branch Office Treatment Roster

Treated offices			Never-treated offices		
Office	First attack	N	Office	First attack	N
Dresden	2015-01	20	Bad Fallingbostel	—	3,208
Hamburg	2015-02	1,157	Bamberg	—	1,098
Nostorf-Horst	2015-02	1,250	Bayreuth	—	494
Manching	2015-04	139	Bramsche	—	4,109
Heidelberg	2015-05	1,479	Buedingen	—	1,624
Berlin Ankunfts.	2015-06	4,987	Deggendorf	—	1,399
Karlsruhe	2015-07	3,833	Flgh. Frankfurt	—	63
Chemnitz	2015-07	3,103	Jena-Hermsdorf	—	1,458
Halberstadt	2015-08	3,090	Moenchengladbach	—	4,177
Leipzig	2015-08	648	Neumuenster	—	3,605
Bremen	2015-09	1,377	Neustadt	—	872
Dortmund	2015-09	4,712	Nuernberg	—	9,530
Unna	2015-09	997	Oldenburg	—	1,306
Essen	2015-09	1,613	Regensburg	—	1,391
Bochum	2015-09	127	Schweinfurt	—	1,353
Friedland	2015-10	2,046	Sigmaringen	—	1,010
Eisenhuettenstadt	2015-10	3,329	Trier	—	4,887
Augsburg	2015-12	790	Zirndorf	—	4,859
Braunschweig I	2016-01	3,345			
Ellwangen	2016-01	2,435			
Muenchen	2016-03	5,507			
Giessen	2016-05	9,114			
Bielefeld	2016-05	7,289			
Duesseldorf	2016-06	4,793			
Bonn	2016-06	200			
Koeln-Bonn	2016-06	1,896			
Suhl	2017-01	1,521			
Freiburg	2017-04	599			
Lebach	2017-06	1,381			

Notes: The 43 immigration office branch offices in the estimation sample (27 treated, 16 never-treated). Treated offices experienced at least one arson attack within 25 km; first attack month shown. Never-treated offices had no nearby arson attack during 2015–2021. N is the number of estimation-sample applicants at each office, pooled across pre- and post-treatment filing periods. The earliest treatment cohorts (first attacked in early 2015) contribute fewer pre-treatment observations. The Berlin arrival center, the Bonn branch, the Frankfurt Airport transit office, and the Nuremberg headquarters are excluded by the regular-application filter (Table A.3). ^a Pre-treatment filing cell below the disclosure threshold (fewer than three persons) is suppressed and excluded from the total. Source: AZR Research Dataset 2023; parliamentary interpellations (arson events).




Figure A.1: Decision Hazard and Cumulative Incidence in the Estimation Sample

Notes: Panel (a) plots the discrete-time hazard rate: the share of the at-risk population receiving a decision in each month after application, separately for positive decisions (refugee status under AsylG or GFK, subsidiary protection, or deportation ban) and rejections. Panel (b) plots the corresponding cumulative incidence: the cumulative share of the initial population that has received each decision type by month t . Month 24 is a right-censored catch-all (applications with Dublin transfers, cancellations, or pending status) and is excluded from the plotted range. Data pooled over application years 2015–2018. Source: AZR Research Dataset 2023.




Figure A.2: Duration from Application to Interview

Notes: Mean months from asylum application filing to personal interview (*Anhörung*), by quarter. Lines show quarterly averages for the overall caseload and for Syria. Shaded region marks the analysis sample (2015–2018). Gaps indicate quarters not covered by the parliamentary inquiry source. Source: Federal parliamentary inquiry data on immigration office processing times, 2013–2018.




Figure A.3: Duration from Interview to Decision

Notes: Mean months from personal interview (*Anhörung*) to final branch-office decision, by quarter. Lines show quarterly averages for the overall caseload and for Syria. Shaded region marks the analysis sample (2015–2018). Gaps indicate quarters not covered by the parliamentary inquiry source. Source: Federal parliamentary inquiry data on immigration office processing times, 2013–2018.

B Identification

B.1 Identification Threats

This subsection provides detailed discussion of each identification concern introduced in Section 5.1.

Allocation of refugees. The identification strategy leverages the quasi-random allocation of refugees across German counties. While certain counties specialize in processing applications from specific countries of origin, these specializations remain largely stable over time. The key identifying assumption is that the distribution of refugees across German regions does not systematically change in correlation with local anti-immigrant sentiment. Even if counties attempted to select applicants based on local attitudes, the specifications control for all refugee characteristics available for selection, including country of origin, gender, age, and family status. Figure B.1 provides direct evidence on this point: predicted approval rates, computed from cell means of nationality, calendar month, gender, and minor status, show no shift around treatment timing.

Selective attrition. The CRF does not retain records of individuals who naturalize. If naturalization rates differ between treated and control counties, the 2023 extract would disproportionately lose approved cases from the group with higher naturalization, biasing observed approval rates downward for that group.¹¹ The structural timeline limits this concern in either direction. The standard naturalization path requires 6-8 years of lawful residence, so applicants recognized in 2015 would satisfy the residence requirement in 2021 at the earliest. The results are robust to controlling for place of residence, meaning that for this channel to drive the results, it would need to operate through the processing location rather than residence. Table B.1 tests this directly by comparing treated and control counties on two administrative margins: register attrition, defined as the log ratio of protection-title holder counts between the 2021 and 2023 CRF extracts,¹² and cumulative naturalizations as a share of the county’s foreign population. Neither measure differs significantly between treated and control counties, whether considering all counties or restricting to those hosting an immigration office. The latter comparison more closely mirrors the unit of treatment variation in the main analysis.

¹¹For attrition to generate the estimated negative treatment effect, treatment would need to *increase* naturalization among affected counties—removing more approved cases from the treated group and depressing its observed approval rate. The reverse channel, in which treatment lowers approval rates and thereby reduces naturalization eligibility, would attenuate the estimate toward zero.

¹²The main analysis uses the AZR Research Dataset 2023 (Gleiser and Salas Escobar, 2024). An earlier extract, the AZR Research Dataset 2021 (BAMF-Forschungsdatenzentrum, 2023), drawn before most applicants in the 2015–2018 cohort could have satisfied the residence requirement for naturalization, provides a pre-naturalization benchmark. This extract contains fewer variables and could not be used for the primary analysis, but suffices for comparing holder counts across cohorts. The ratio of holder counts between the two extracts captures all sources of register attrition.

Exogeneity of attack timing. The identifying assumption requires that the timing of first arson attacks near immigration offices is uncorrelated with unobserved determinants of approval rates, conditional on the fixed effects. Arson attacks are not randomly assigned: they cluster in areas with rising anti-immigrant sentiment. The conceptual framework treats arson attacks as salient signals through which officers learn about local anti-immigrant sentiment. The identification concern is not that sentiment exists, but that offices may already be responding to the same underlying sentiment through other channels, such as colleague conversations, local media tone, or community attitudes, before the attack occurs. If so, the attack would be a symptom of sentiment that has already influenced decisions rather than an independent informational shock. Three features of the design mitigate this concern. First, the calendar month $\times$ nationality $\times$ duration fixed effects absorb nationwide sentiment shifts. Second, branch office fixed effects absorb time-invariant local attitudes. Third, residence county fixed effects absorb cross-county heterogeneity in applicants' locations, and the robustness specification that interacts residence county with calendar month (Table C.2) additionally absorbs time-varying shocks at the county level; the estimate is, if anything, larger under that specification. What remains unabsorbed is time-varying sentiment *specific to the branch office's locality* that correlates with attack timing. I cannot fully rule out this threat. The event study's pre-treatment coefficients provide indirect evidence: in the preferred specification, none of the lead coefficients is statistically significant (Figure 3).

Spillover. The event study design assumes that arson attacks near one branch office do not affect decisions at other offices. Two types of spillover are relevant. First, geographic spillover: news coverage of an attack may affect applicants' behavior or legal-support availability near the attack site, independent of the processing office. Residence county fixed effects absorb level differences across applicants' locations, the residence county $\times$ calendar month robustness specification (Table C.2) absorbs time-varying local shocks, and the restriction to applicants whose residence county differs from their processing office county directly probes this channel. Second, office-to-office spillover: staff at untreated offices may respond to attacks at distant offices through internal immigration office communication or national media coverage. If such spillover depresses approval rates at control offices, the treatment effect is attenuated; if it elevates them (e.g., through a compensatory response), the estimate is inflated. Under non-uniform spillover, the dynamic coefficients can be distorted in either direction, and attenuation is not guaranteed. Results are robust to alternative control groups—never-treated offices only, and the full not-yet-treated pool—which would be differentially affected if internal spillover were pervasive. Nevertheless, spillover biases cannot fully be ruled out.

Institutional responses. Immigration office leadership may have responded to arson attacks with internal directives, changes in oversight intensity, or personnel reallocation—

channels distinct from individual officer discretion. If such institutional responses systematically lowered approval rates at affected offices, the estimated effect would conflate individual behavioral change with top-down institutional reactions. As discussed in Section 5.1, the federal structure of asylum adjudication limits any branch-office-level managerial reaction to procedural adjustments rather than substantive policy change. Without access to internal communications or personnel records, however, I cannot rule out procedural responses such as shifting case assignments or intensifying supervisory review at affected offices. Following the discovery of procedural irregularities at the Bremen branch office in mid-2018, the immigration office introduced systematic monitoring of branch-office approval rates against a country-composition-adjusted reference rate, with deviations exceeding 10 percentage points triggering central review.¹³ This reform falls at the trailing edge of the sample period and does not drive the main results—the through-2017 robustness specification predates it entirely—but its introduction confirms that headquarters recognized cross-office variation in adjudication as a quality concern.

Legal representation. Arson attacks may trigger increased legal aid or NGO mobilization near affected areas, which could independently affect approval rates by improving case preparation. Since applicants reside in different counties from their processing offices, this channel would need to operate through the applicant’s residence county rather than the office’s locality. Residence county fixed effects absorb time-invariant differences in legal support across applicant locations, and the residence county $\times$ calendar month robustness specification (Table C.2) additionally absorbs time-varying differences, limiting this channel to within-county-month variation in legal mobilization.

B.2 Balance Tests

The balance test in Figure B.1 checks whether applicant composition, as summarized by predicted approval rates, shifts systematically around the timing of a first nearby arson attack. Predicted approval rates are computed from cell means (nationality $\times$ calendar month $\times$ gender $\times$ minor status) using either pre-treatment periods or the never-treated sample, which summarize an applicant’s expected approval probability based on observable characteristics alone. I then regress each predicted rate on the same event-study specification used for the main outcome (Equation 6), restricting to the first observation per application. If treatment status is uncorrelated with applicant composition, the event-time coefficients should be statistically indistinguishable from zero throughout the event window. Significant pre- or post-treatment movements in predicted approval rates would indicate compositional change that the calendar month $\times$ duration $\times$ nationality fixed effects in the main specification are absorbing. Both variants—using pre-treatment

¹³The reference rate weights each branch office’s nationality mix by the national approval rate per nationality, producing a counterfactual mean against which the office’s actual rate is compared (source: parliamentary inquiries BT-Drs. 19/6786 and 20/14882).

Table B.1: Balance: treated vs. control counties

	Treated	Control	Diff.	p
<i>Panel A: All counties</i>				
Register attrition ($\ln N^{2023}/N^{2021}$)	-0.070 (0.022)	-0.029 (0.029)	-0.041 (0.036)	0.261
Naturalizations 2020–2023 (% of foreign pop.)	5.68 (0.14)	6.08 (0.21)	-0.40 (0.25)	0.120
N	227	140		
<i>Panel B: Immigration office counties</i>				
Register attrition ($\ln N^{2023}/N^{2021}$)	0.046 (0.083)	0.105 (0.099)	-0.059 (0.129)	0.649
Naturalizations 2020–2023 (% of foreign pop.)	5.91 (0.35)	5.74 (0.63)	0.17 (0.72)	0.820
N	33	16		

Notes: Treated counties experienced at least one arson attack within 25 km. Control counties are never-treated. Register attrition defined as $\ln(N_c^{2023}/N_c^{2021})$ where N_c^t is the total number of protection-title holders in county c across the 2015–2018 recognition cohorts in the CRF extract of year t . Naturalizations are cumulative 2020–2023 counts as a percentage of the county’s 2020 foreign population (Destatis Einbürgerungsstatistik). Standard errors of the mean in parentheses. Immigration office counties are those hosting a branch office of the Federal Office for Migration and Refugees. Sample restricted to counties observed in both datasets.

and never-treated cell means—show no evidence of compositional change.




Figure B.1: Balance Tests: Predicted Approval Rates around Treatment

Notes: Event-study coefficients from regressions of predicted approval rates on event-time dummies (stacking estimator, Equation 6), relative to the month before treatment ($k = -1$). Predicted rates are computed from cell means (nationality $\times$ calendar month $\times$ gender $\times$ minor status) using pre-treatment periods (a) or the never-treated sample (b). Sample restricted to the first observation per application. All specifications include branch office $\times$ sub-experiment, residence county $\times$ sub-experiment, and calendar month $\times$ duration $\times$ nationality $\times$ sub-experiment fixed effects. Standard errors clustered at the branch office level. Vertical bars show 95% confidence intervals. Source: AZR Research Dataset 2023; parliamentary interpellations (arson events).

C Additional Results

C.1 Heterogeneous Effects

The effect varies across applicant groups (Table C.1). Geographically, the estimate is identified from West Germany; a separate East Germany split is uninformative because nearly all eastern offices experienced their first nearby attack before or early in the sample window, leaving almost no never-treated comparison group within the East. Areas with below-median AfD vote share in the 2013 federal election drive the result, while high-AfD areas show a smaller, insignificant response—consistent with officers in areas where anti-immigrant sentiment is already high having already adjusted their baseline behavior. Attacks classified as aggravated arson are associated with larger effects than simple arson, though the difference between the two estimates is imprecise. Among demographic groups, the decline is concentrated among male applicants, while the female estimate is smaller and statistically insignificant. Applicants from Muslim-majority origin countries show a larger point estimate than other applicants, but the estimate falls short of conventional significance. The expedited-procedure split is too imprecise to be informative: the standard error for the expedited group is roughly five times that of the main estimate.

C.2 Robustness

Table C.1: Heterogeneous effects: Post-period DiD summary

Subgroup	Post $\times$ Treated	SE	N	Persons
<i>Geography</i>				
West Germany	-0.008***	(0.003)	2,444,981	82,073
<i>Political context</i>				
Low AfD vote share (2013)	-0.010**	(0.004)	2,458,082	63,171
High AfD vote share (2013)	-0.003	(0.003)	2,457,838	66,083
<i>Attack severity</i>				
Aggravated arson	-0.008**	(0.003)	2,515,796	71,521
Simple arson	-0.004	(0.004)	2,400,124	57,733
<i>Demographics</i>				
Male	-0.008***	(0.003)	2,547,045	76,217
Female	-0.003	(0.003)	2,368,875	53,037
Muslim-majority origin	-0.008	(0.005)	2,438,292	65,268
Other origin	0.000	(0.001)	2,477,628	63,986
<i>Procedure type</i>				
Expedited procedure	-0.006	(0.014)	2,324,222	49,986
Regular procedure	-0.001	(0.001)	2,591,698	79,268

Notes: Summary of post-period difference-in-differences estimates across heterogeneity dimensions. Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Each row corresponds to a subgroup estimated from a separate stacking regression with a single post-treatment indicator (never-treated comparison). Splits classify observations by the indicated attribute while retaining the full never-treated comparison group; the West Germany row instead restricts the sample to western states. N denotes person-months in the stacked panel; Persons denotes unique applicants. All specifications include branch office $\times$ sub-experiment, residence county $\times$ sub-experiment, and calendar month $\times$ duration $\times$ nationality $\times$ sub-experiment fixed effects, and demographic controls. Standard errors clustered at the branch office level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table C.2: Robustness: Fixed-Effect Specifications

	(1)	(2)	(3)	(4)	(5)
Post $\times$ Treated	-0.005** (0.002)	-0.005* (0.003)	-0.006** (0.003)	-0.009** (0.004)	-0.009** (0.004)
Mean dep. var.	0.026	0.026	0.026	0.026	0.026
Observations	2,682,401	2,682,401	2,682,401	2,682,401	2,682,401
Persons	95,320	95,320	95,320	95,320	95,320
Office $\times$ sub-exp.	Yes	Yes	Yes	Yes	Yes
Res. county $\times$ sub-exp.	Yes	Yes	Yes		Yes
Duration $\times$ sub-exp.	Yes	Yes			
Month $\times$ sub-exp.	Yes				
Nationality $\times$ sub-exp.	Yes				
Month $\times$ nat. $\times$ sub-exp.		Yes			
Month $\times$ dur. $\times$ nat. $\times$ sub-exp.			Yes	Yes	Yes
Res. county $\times$ month $\times$ sub-exp.				Yes	
Office state $\times$ month $\times$ sub-exp.					Yes

Notes: Each column reports the collapsed post-treatment coefficient from the stacking estimator (Equation 5) under a different fixed-effect structure, indicated in the lower panel. Dependent variable: any protection (never-treated comparison). Column (3) corresponds to the main specification in Table 2. All specifications include demographic controls (gender, family status, birth year, minor status). Standard errors clustered at the branch office level in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table C.3: Robustness: Maximum Observation Duration

	(1) 9 months	(2) 6 months	(3) 3 months
Post $\times$ Treated	-0.008*** (0.003)	-0.010** (0.004)	-0.010* (0.005)
Mean dep. var.	0.025	0.026	0.024
Observations	2,801,863	2,217,006	1,423,076
Persons	95,938	95,938	95,938

Notes: Each column reports the collapsed post-treatment coefficient from the stacking estimator (Equation 5) restricting the panel to the indicated number of months after application. Dependent variable: any protection. These specifications pool the never-treated and not-yet-treated comparison groups. All specifications include branch office $\times$ sub-experiment, residence county $\times$ sub-experiment, and calendar month $\times$ duration $\times$ nationality $\times$ sub-experiment fixed effects, as well as demographic controls. Standard errors clustered at the branch office level in parentheses. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.