Political Violence and Bureaucratic Discretion: Evidence from German Asylum Decisions

Yogam Tchokni*

31/05/2026

[Click here for the most recent version]

Abstract

Democratic accountability typically flows downward: voters pressure politicians, who direct agencies. I test whether policy also responds from below—through the discretionary decisions of the bureaucrats who implement it. I study German asylum decisions, exploiting arson attacks against refugee accommodations as salient signals of anti-immigrant hostility. Unlike out-group-perpetrated violence studied previously, in-group attacks against a vulnerable minority activate competing pressures on decision-makers: compliance with hostile demands versus empathy for victims. Using administrative data on 148,405 applicants (2016–2018), immigration offices exposed to a nearby attack show a decline in approval rates, without a change in policy directives. Media-salient attacks produce larger effects and severe attacks show a partial offset, consistent with the empathy channel.

Keywords: Immigration, Political Violence, Asylum, Discrimination, Bureaucratic Discretion

JEL codes: D73, J15, K37

*DIW Berlin, ytchokni@diw.de

1 Introduction

What if bureaucrats take the law into their own hands? Democratic accountability is typically understood as flowing downward: voters pressure politicians, who direct agencies. But policy implementation depends on the people who carry it out. Across courts, welfare offices, and government agencies, public servants exercise substantial discretion (Lipsky, 2010). When implementers possess such power, their decisions may respond not only to demands from policymakers but also to the social and political conditions of the local environment in which they operate. Policy may therefore be responsive to demands from the *bottom-up* in the population beyond traditional *top-down* channels.

I test whether signals of extreme and violent opposition to government policy affect its execution through public servants' use of bureaucratic discretion. I use the setting of German asylum decisions and the wave of political violence against refugees from 2015 to 2018. Applications that were processed after attacks near the immigration office are less likely to be approved.

The German asylum system is well-suited to study this question. First, asylum policy is determined at the federal level, so local variation in approval rates reflects adjudication rather than policy change. Second, applicants are quasi-randomly allocated across Germany through a centralized distribution system, limiting sorting. Third, geographic separation between applicants' residences and the 48 immigration office branch offices that process their claims allows me to separate office-side exposure to local hostility from exposure near applicants' own residences. Arson attacks against refugee accommodations, perpetrated by members of the majority population against a vulnerable minority, serve as salient, localized signals of extreme and violent opposition to existing government policy.

Results are consistent with a *bottom-up* channel in which approval rates comply with hostile demands. A within-office test provides direct evidence for case officer discretion at the personal interview as the site of discretion: the arson effect is concentrated among cases adjudicated through a personal hearing and absent for cases decided via written procedure at the same office during the same period. Because the administrative data do not record the full history of each case, interpreting effects as a change in approval rates requires assuming that treatment does not affect decision timing. A robustness test using aggregate office-level statistics confirms no change in the volume of decisions at treated offices. I also examine two exploratory dimensions. The estimated effect is larger for attacks receiving substantial media coverage, consistent with an information channel identified in previous research (Emeriau, 2024; Spirig, 2023). More severe attacks involving risk to life show a partial offset at longer horizons, aligned with research showing extreme violence can provoke a countervailing sympathetic response toward targets. Survey evidence from the SOEP shows that arson attacks shift public attitudes among residents of immigration office counties in a way consistent with compliance dominating resistance as responses to the political violence. Concerns about immigration rise persistently, paralleling the durable voting shifts Sabet et al. (2023) documents. Concerns about xenophobic hostility spike immediately but dissipate within weeks, making any resistive responses short-lived, consistent with Emeriau (2024) and Freitas-Monteiro and Prömel (2024). The asymmetry provides attitudinal evidence for why the bureaucratic response tilts toward stricter treatment.

A large body of literature has shown that discretion shapes outcomes across all aspects of government, such as legal proceedings (Shayo and Zussman, 2011; Arnold et al., 2018; Eren and Mocan, 2018; Dobbie et al., 2018), policing (Grosjean et al., 2023), welfare allocation (Hemker and Rink, 2017; Maestas et al., 2013; Bolhaar et al., 2020) or immigration (Emeriau, 2023, 2024; Chen et al., 2016; Gundacker et al., 2024).

Researchers have shown that decision-makers respond to external events focusing on three categories: out-group-perpetrated violence (Shayo and Zussman, 2011; Brodeur and Wright, 2019; Gould and Klor, 2016; Emeriau, 2024), emotional shocks (Chen and Loecher, 2025; Emeriau, 2024; Heyes and Saberian, 2019; Danziger et al., 2011), and inflammatory rhetoric (Grosjean et al., 2023; Spirig, 2023). Shayo and Zussman (2011) show that judges exhibit heightened in-group bias following terrorist attacks, that is, violence perpetrated by members of the out-group. Brodeur and Wright (2019) find that the 9/11 attacks reduced asylum approval rates for Muslim applicants in the United States, again responding to out-group violence. For random events, Eren and Mocan (2018) document that unexpected losses in local football games lead to harsher juvenile sentencing. Emeriau (2024) shows that news of Mediterranean shipwrecks increases asylum approval rates in France through empathy. The paper most closely connects to literature on the effects of inflammatory rhetoric. Grosjean et al. (2023) show that inflammatory rhetoric at Trump campaign rallies increased racial profiling, and Spirig (2023) shows that media reporting of anti-immigrant sentiment and crime reduces judicial approval rates.

Recent work on far-right political violence documents effects on voting behavior and political rhetoric (Sabet et al., 2023; Jomni and Zakharov, 2026), establishing that such violence can increase support for far-right parties, thereby affecting policy through pressure on politicians. On the other hand, Freitas-Monteiro and Prömel (2024) show that large far-right demonstrations in the same setting increased concern about hostility toward foreigners and pro-refugee behaviour. How far-right violence shapes bureaucratic decision-making, and in which direction, is an open question. The most closely related studies are Spirig (2023), who shows that media salience of immigration reduces judicial approval rates in Switzerland, and Gundacker et al. (2024), who find lower approval probabilities for refugees living in immigration-averse regions of Germany.

This paper makes three contributions. First, I provide evidence on a qualitatively distinct hostility signal: *in-group* violence against a vulnerable minority. Unlike out-group-perpetrated attacks, the predicted direction of the bureaucratic response is ambiguous. Such violence could normalize hostility (Bursztyn et al., 2020; Sabet et al., 2023) and serve as a signal of extreme opposition to current policy, demanding stricter bureaucratic treatment. On the other hand, violent or extremist protests may also discredit political positions (Bursztyn et al., 2020; Chenoweth and Stephan, 2011; Feinberg et al., 2020; Freitas-Monteiro and Prömel, 2024), and create empathy for the targeted group. Existing work focused on out-group-perpetrated violence can be interpreted through rational decision-making given updated threat perceptions. In-group violence against a vulnerable out-group offers no such mechanism as the targeted group poses no greater threat after the attack. Instead, it generates competing pressures on decision-makers: compliance with hostile demands versus empathy driven leniency. I show that the former wins out.

Second, the geographic separation between processing offices and applicant residences enables a design that isolates the office-side channel from applicant-side confounders such as access to legal aid and differential behavior that could otherwise also be impacted by arson attacks. The closest existing study is Gundacker et al. (2024), who find lower approval rates for individuals residing in immigration-averse regions.

Third, I provide the first analysis of asylum adjudication using individual-level administrative data from Germany, the largest recipient of asylum applications in Europe. Between 2015 and 2018, Germany received approximately 1.5 million first-time applications, accounting for 40 percent of all applications filed across Europe. Previous studies using administrative data have focused on countries with substantially smaller caseloads such as France (9 percent) (Emeriau, 2024) and Switzerland (3 percent) (Spirig, 2023).

Section 2 describes the German asylum system. Section 3 develops the conceptual framework. Section 4 presents the data, Section 5 the empirical strategy, and Section 6 the results. Section 7 concludes.

2 Institutional Setting

This section describes the institutional features of the German asylum system.

2.1 Allocation System

Germany administers refugee arrivals through a centralized allocation system. Upon arrival, asylum seekers register at reception centers and are allocated to federal states via the EASY system (*Erstverteilung der Asylbegehrenden*), which follows the Königsteiner Schlüssel, a formula weighted by each state's population and tax revenue. States then distribute individuals to municipalities based on housing availability. Allocation decisions are largely independent of individual refugee characteristics (Jaschke, Sardoschau, and Tabellini, Jaschke et al.; Aksoy et al., 2023).

Following municipal assignment, applicants are directed to one of 48 immigration office branch offices for processing, often located in different counties than applicants' residences. This geographic wedge is what allows me to separate office-side exposure from applicant-side exposure: an arson attack near an immigration office reaches that office's staff and local environment, but not applicants residing in distant counties who happen to be processed there.

2.2 Decision Process

Each applicant undergoes a personal interview conducted by an asylum officer. The interview forms the basis for assessing eligibility for protection. Officers have substantial discretion: they control the line of questioning, assess the credibility of testimony, and ultimately determine both whether to grant protection and what level of protection to award.

Applicants may receive one of several outcomes: full refugee status under the Geneva Convention (three-year permit, immediate family reunification), subsidiary protection (one-year permit, restricted reunification), a deportation ban, or rejection. Rejected applicants may appeal to administrative courts. See Appendix Table A.2 for detailed descriptions of protection levels. Figure 2 illustrates the observability of each outcome type under the 12-month restriction used in the empirical design.

2.3 Discretion

The interview serves as the primary opportunity for decision-makers to assess the credibility and substance of an asylum claim. During this meeting, officers evaluate the applicant's testimony, probe for inconsistencies, and form judgments about protection eligibility. While applications are filed and administrative processing occurs over extended periods, the interview represents the concentrated moment when subjective evaluation takes place. Events occurring close to this critical decision point are therefore most likely to influence officer judgments, whether through heightened threat perceptions following local violence or through increased empathy following humanitarian tragedies. Following Emeriau (2024), who uses interview dates to identify the timing of discretionary effects on asylum decisions in France, I treat the interview as the critical moment for adjudication.

The degree of discretion varies by applicant origin. During 2016–2018, Syrian applicants were rarely rejected outright, though officers exercised discretion over whether to grant full refugee status or the less favorable subsidiary protection.[1] For applicants from other origins, officer discretion extends to the binary approval decision as well. This variation in the scope of discretion generates a testable prediction: if attacks operate through officer

[1] In 2016, 167 of 266,866 Syrian applications were rejected outright; of the remainder, 165,754 received full Geneva-Convention protection and the balance received subsidiary protection, which restricts family reunification.

discretion at the interview, effects should concentrate among nationalities where officers have more latitude over outcomes.

The scope for discretion was particularly wide at the start of the study period due to weak internal oversight. The immigration office's central quality-assurance unit reviewed approximately 1% of decisions in 2015, and the dual-review requirement for individual decisions was made formally binding by central directive only in September 2017 (see Appendix A for details). Until at least late 2017, officers operated with minimal systematic review of their individual decisions.

3 Conceptual Framework

I conceptualize arson attacks against refugee accommodations as salient events that increase the local visibility of anti-immigrant sentiment and social dissatisfaction. Such attacks signal that segments of the population hold strong anti-refugee views—views intense enough to motivate violent action. The existing literature on how salient events affect bureaucratic and judicial decision-makers suggests several pathways through which such attacks may alter officer behavior. The predicted direction is ambiguous ex ante, as the same event can produce both compliance with and resistance to political violence.

Shift in own beliefs. Arson attacks signal that social cohesion in the local community is under strain. Officers who learn of violence against refugee accommodations may perceive their surroundings as destabilized. This perceived disruption can harden attitudes toward the group associated with the tension, even when that group is itself the target. Sabet et al. (2023) show that successful right-wing attacks in Germany shift individual attitudes rightward and increase anti-immigration concern, even among previously unaffiliated individuals. Salient events do not need to change beliefs from nothing but can instead activate pre-dispositions that already exist. Grosjean et al. (2023) show that Trump campaign rallies increased racial bias in policing, with effects concentrated among officers whose baseline behavior was already more biased. In the immigration office context, rapid hiring during 2015–2016 brought officers with heterogeneous predispositions into the same offices. An arson attack nearby may both shift attitudes on the margin and activate restrictive tendencies among officers already predisposed—two sides of the same process, requiring neither strategic calculation nor conscious awareness.

Shift in perceived expectations. A distinct but complementary channel operates through perceived expectations rather than personal beliefs. Attacks may not change what officers think about refugees, but instead what officers believe their environment expects of them. Spirig (2023) shows that higher asylum salience in Swiss media leads to more restrictive judicial decisions on asylum appeals and is not restricted to judges of any particular ideology. Officers may calibrate decisions toward what they perceive as the locally appropriate position, whether through internalized salience effects or through social pressure from colleagues and the community atmosphere (Lipsky, 2010). This channel does not require officers to personally hold anti-refugee views. The institutional context of 2016–2017 amplifies this pathway: the immigration office was rapidly expanding its workforce, and new hires on probationary periods may have been especially susceptible to signals about local expectations. Emeriau (2023) shows that new hires are often more biased than experienced ones.

Protective action. Arson attacks physically endanger refugees in their accommodations. Given that positive asylum decisions allow refugees to choose their place of residence more freely, officers confronted with physical endangerment may view applicants more favorably. Emeriau (2024) shows that news of Mediterranean shipwrecks increased asylum approval rates in France by 4.4 percentage points when featured on prime-time television. However, this empathy effect was extremely short-lived, disappearing within one day of the event. Given that arson attacks occur within a context of social tension rather than humanitarian disaster, it is unclear whether they trigger the same victim framing.

Backlash against extremism. A broader literature on responses to political violence shows that extreme tactics frequently activate moral outrage and sympathy for targets (Wasow, 2020; Feinberg et al., 2020; Chenoweth and Stephan, 2011). In the German context, Freitas-Monteiro and Prömel (2024) document how far-right demonstrations increase concern about hostility toward foreigners, raise intentions to donate to refugee aid, and benefit left-wing parties. Officers may feel revulsion toward the attackers and compensate by exercising discretion more favorably.

The framework does not provide a clear prediction for the sign of the net effect. The resistance pathways require specific conditions—victim framing and immediate temporal proximity for empathy, extreme visible violence for backlash—that may not hold for arson attacks in this context. Attack severity provides one dimension of leverage: if empathy or backlash operates, aggravated arson that endangers human life should show a different pattern than minor property damage. Whether attacks receive media coverage may matter for the salience channel. I return to these distinctions when interpreting the heterogeneity results in Section 6.

4 Data and Measurement

Figure 1 displays the geographic and temporal distribution of arson attacks on refugee accommodations across German counties, which constitute the treatment events in the empirical design.

4.1 Attack data

A key source of administrative data on anti-refugee violence comes from parliamentary interpellations submitted to the German federal government. These interpellations are formal inquiries that require the government to provide detailed information on specific topics, in this case regarding attacks on refugee accommodations and asylum seekers.

The federal government compiles responses using data from the Federal Criminal Police Office through the Criminal Police Reporting Service for Politically Motivated Crime. These reports provide incident-level data on attacks on asylum accommodations (both occupied facilities and those under construction), attacks on asylum seekers in public spaces, and far-right protests at or near refugee facilities. Each incident is recorded with the date, location (city and state), and the specific criminal offense.

I use quarterly interpellations covering 2015–2021, which provide the most systematic and comprehensive administrative record of anti-refugee violence during this period. The full 2015–2021 window is used to determine the timing of each office's first nearby attack, ensuring that offices attacked before the 2016–2018 estimation sample are correctly classified as already treated. Figure 1 displays the geographic distribution of these attacks across German counties over time.

The data have two advantages over media-based compilations. First, they are based on official police reports rather than media coverage, which reduces concerns about selective reporting. Second, they provide standardized legal classifications of offenses, enabling consistent severity coding across incidents and years.

The data classify arson incidents according to the German Criminal Code. I distinguish between simple arson and aggravated arson, where the latter involves endangering human life, occupied buildings, or serious bodily harm.[2] I construct an *aggravated arson* indicator for attacks classified under the aggravated categories, representing offenses with higher potential risk to human life. This distinction allows me to examine whether the severity of attacks moderates their effect on asylum decisions.

4.2 Application data

My main data source is the Central Register of Foreigners (CRF), a nationwide administrative register that is centrally managed by a federal office. It contains personal data, nationality, residency status, and other legal matters such as asylum applications, for all non-German citizens residing in or having resided in the country. The immigration office provides a 20% random sample for research purposes (Gleiser and Salas Escobar, 2024). The sample was drawn from the CRF on 31 December 2023.

The data contain information on place of residence, residence status, and other personal

[2]See Appendix Table A.1 for the full classification of criminal offenses.

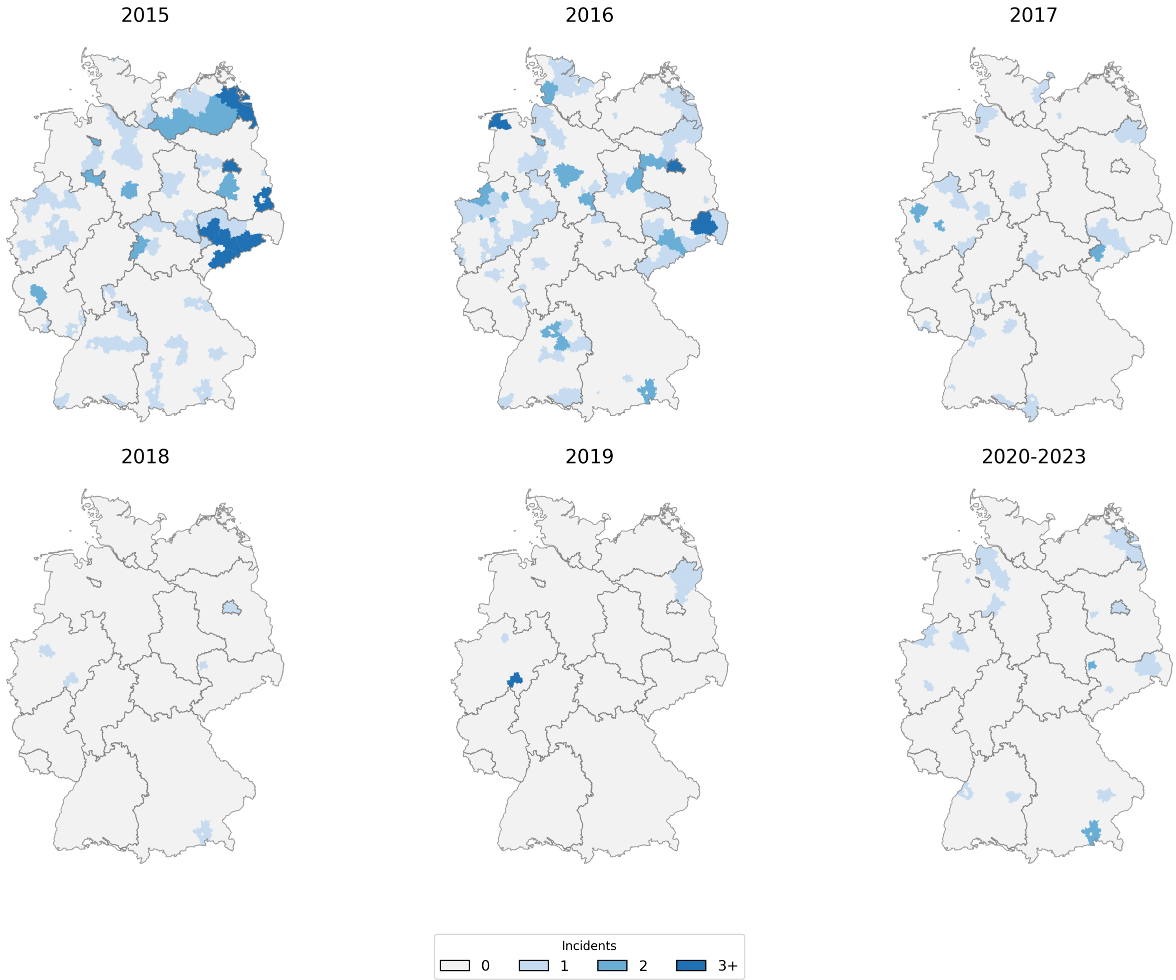

Figure 1: Arson Attacks on Refugee Accommodations by County and Year
Notes: Data from parliamentary interpellations to the German federal government, compiled from Federal Criminal Police Office records. Counties are shaded by the number of arson attacks on refugee accommodations recorded in each year. The 2020–2023 panel aggregates the post-sample period for comparison.

characteristics, with records entered by local municipalities, immigration offices, the police, and other authorities. For refugees, the residence-status history tracks the progression of the asylum application. Socio-demographic variables include age, gender, country of origin, education, and previous employment. Coverage is limited to non-citizens at the time of sample draw, since foreigners who naturalize are deleted from the CRF upon obtaining German citizenship. Appendix B.1 discusses whether selective sample attrition threatens identification.

I restrict the estimation sample to person-month observations in calendar years 2016–2018. During 2015, the unprecedented volume of asylum applications led the immigration office to adopt simplified approval procedures for several nationalities, including expedited written decisions without individual interviews.[3] Since the bottom-up interpretation relies

[3]From 2014 through late 2015, the immigration office granted refugee status to Syrian and Eritrean nationals and members of Iraqi religious minorities through a written fast-track procedure on the basis of a questionnaire alone.

on adjudication exercised at the personal interview, including 2015 filings, most of which bypassed the discretionary stage, would dilute the sample with outcomes unaffected by local conditions. I therefore add an explicit filing-date floor of 1 January 2016: only applications filed on or after this date enter the panel. The filing-date restriction removes applicants filed in late 2015 who, absent the restriction, would contribute 2016 person-months under the calendar-year filter alone.

The resulting sample restrictions and their effect on applicant counts are summarized in Table A.3. Treatment timing is computed over the full 2015–2021 interpellation window, so offices that experienced their first nearby arson attack in 2015 enter the sample as already treated.

One key limitation of the CRF data is that it only contains legally binding changes to the asylum status. Therefore asylum status is only observed once the legal appeals process has been exhausted (see Figure 2). Interview dates, initial decisions and their timing can therefore not be directly inferred from the data, so the treatment-relevant window cannot be precisely identified. However, the duration between application and final status provides information on the initial decision and motivates the use of a duration model. The following section describes how these features are leveraged.

(a) Case A: Observed Approval

(b) Case B: Observed Rejection

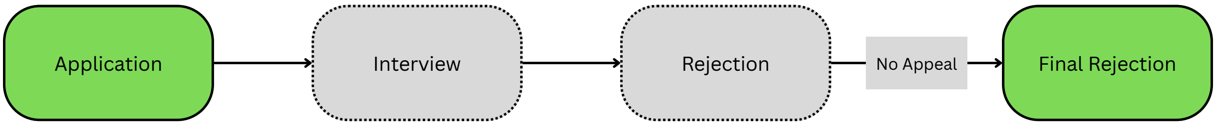

(c) Case C: Appeals Outcome

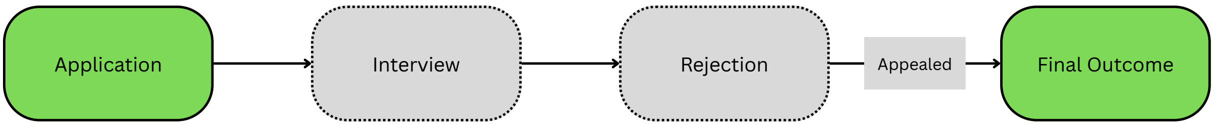

Figure 2: Data Observability in Asylum Decision Processes

Notes: Schematic illustration of three types of asylum application outcomes and their observability under the 12-month restriction. Green shading indicates stages observed in the data; grey shading indicates stages that are unobserved. Case A: positive decision within 12 months (observed). Case B: negative decision within 12 months (observed as zero in outcome variable). Case C: appeal process extending beyond 12 months (applicant exits sample as censored). The 12-month window is designed to capture first-instance determinations while excluding most appeals.

5 Empirical Strategy

5.1 Identification Strategy

The CRF records only the final legally binding asylum status, observed after any appeals have concluded. Approved applications resolve at the time of the first-instance decision, but rejected applications may be appealed and appear in the data only much later. Using the date of the final outcome to define pre- and post-treatment periods would therefore conflate the *content* of decisions with their *timing*: a post-attack decline in observed positive outcomes could reflect either fewer approvals or a shift in when cases resolve.

To separate these two channels, I define the outcome as the monthly approval indicator

$$Y_{it} = \mathbf{1}(\text{applicant } i \text{ receives a positive decision in month } t) \tag{1}$$

with $Y_{it} = 0$ for all months if no approval is recorded within 12 months of application, regardless of whether the applicant was rejected, is still pending, or has entered the appeals process. This creates a person × month panel in which each applicant contributes observations from the month of application until either a positive decision or the 12-month censoring point.

Because approvals are not appealed, any first-instance approval is observed at its true decision date, even though rejections are not directly observed at the correct time. Among still-pending applicants, the expected approval indicator decomposes as:

$$\mathbb{E}[Y_{it}] = \mathbb{E}\Big[\underbrace{\lambda_{it}}_{\substack{\text{decision}\\\text{rendering}}} \cdot \underbrace{Z_i}_{\substack{\text{approval}\\\text{outcome}}}\Big] \tag{2}$$

where λ_{it} is the probability that a first-instance decision of any kind is rendered for applicant i in month t among still-pending applicants, and $Z_i \in \{0, 1\}$ indicates whether the decision is positive. The treatment effect of interest is the change in $\mathbb{E}[Z_i]$. Recovering this from the observed left-hand side requires assumptions that discipline λ_{it} and the composition of the at-risk population. The design rests on three assumptions.

Assumption ID.1 (Exogenous applicant composition). Conditional on place of residence, the allocation of asylum applicants across processing branch offices is orthogonal to treatment status. For applications submitted prior to treatment, office assignment is predetermined and plausibly uncorrelated with treatment. For applications submitted after treatment, I assume that the flow of new filings to each office does not respond to nearby arson attacks. This is plausible because applicants are assigned to offices administratively and cannot choose their processing location. Figure B.1 tests this assumption by showing that predicted approval rates, constructed from pre-treatment cell means, do not shift around treatment timing.

Assumption ID.2 (Appeals fall outside the observation window). The combined first-instance processing time and appeals duration exceeds the 12-month observation window, so the only positive outcomes observed within that window are first-instance approvals: no successful appeal contaminates Y_{it}. This is empirically supported by institutional processing times: average asylum court proceedings took 12.5 months in 2018, on top of several months of first-instance immigration office processing.[4]

Assumption ID.3 (Content, not timing). Exposure to a nearby arson attack shifts the expected approval outcome $\mathbb{E}[Z_i]$ but not the decision-rendering hazard λ_{it}. Let D indicate whether a branch office has been exposed to a nearby arson attack. ID.3 requires that, within calendar month × duration × nationality cells:

$$\mathbb{E}[\lambda_{it} \mid D = 1, t, \tau, n] \; = \; \mathbb{E}[\lambda_{it} \mid D = 0, t, \tau, n] \tag{3}$$

so that λ_{it} is mean-independent of treatment status and any difference in $\mathbb{E}[Y_{it}]$ across treatment arms isolates the shift in $\mathbb{E}[Z_i]$, not in the timing of decisions. I test this assumption directly using office-level aggregate statistics in Section 6.1.1 and find no evidence that the decision-rendering hazard changed at treated offices.

Estimand. Under Assumptions ID.1–ID.3, the difference-in-differences coefficient identifies:

$$\beta \; = \; \mathbb{E}[\lambda] \; \cdot \; \big(\mathbb{E}[Z \mid D = 1] \; - \; \mathbb{E}[Z \mid D = 0]\big) \tag{4}$$

the shift in the approval content of decisions at exposed offices, scaled by the average decision-rendering hazard $\mathbb{E}[\lambda]$. The factorization from Equation 2 to 4 requires that λ_{it} and Z_i are conditionally independent given treatment status and cell membership: ID.3 ensures λ_{it} does not vary with D, and ID.1 ensures the composition of Z_i does not vary with D within cells, so the expectation of their product separates. β therefore isolates the change in approval rates for treated applications.

Treatment coding. A branch office enters the post-first-attack regime upon the first arson attack occurring within a 25 km radius. Subsequent attacks during the sample period are not coded separately, so the post-period coefficients reflect the joint effect of the first attack and any subsequent nearby incidents at the same office. This simplifies the event-study to a single cutoff per office but means the design does not, by itself, decompose the initial shock from cumulative exposure. As robustness checks, I adjust the treatment coding to different distance thresholds.

From office-level estimate to officer discretion. The office-level estimate could reflect three channels beyond individual caseworker discretion: differential managerial direc-

[4]Source: parliamentary interpellations on asylum processing durations (Schwerpunktfragen zur Asylverfahrensdauer, BT-Drs. 20/5709).

tives at treated offices, differential staffing or case-assignment responses, or applicant-side confounders operating within the office's local ecosystem. Because asylum law is federal and immigration office headquarters issues centralized operational guidelines that apply identically to all 48 branch offices, responses that differentially affect offices such as increased oversight would have to be made explicit. The direct channel through which local conditions can enter adjudication is the individual officer's exercise of discretion, in assessing credibility, weighing evidence, and determining protection outcomes. Appendix B.1 provides an extended discussion of institutional responses and other threats.

Threats to identification. Three first-order concerns bear on the design. First, arson attacks cluster in areas with rising anti-immigrant sentiment, raising the possibility that offices are already responding to the same underlying sentiment before the attack occurs. The fixed effects absorb nationwide sentiment shifts, time-invariant local attitudes, and county-level time-varying shocks; what remains unabsorbed is time-varying sentiment specific to the branch office's locality. The event study examines whether pre-treatment coefficients reveal differential trends at treated offices; the four pre-treatment coefficients are individually and jointly insignificant, though with a five-month pre-window the test has limited power to detect violations that would meaningfully bias the post-treatment estimate (Roth, 2022). The design also requires a no-anticipation assumption: officers do not adjust decisions in response to attacks that have not yet occurred. Given that arson attacks are criminal acts without public advance notice, anticipation effects are unlikely.

Second, attacks may affect decisions at distant offices through news coverage or internal communication, violating the stable unit treatment value assumption. Two types of spillover are relevant: geographic spillover via news coverage affecting applicants' locations, and office-to-office spillover via internal immigration office communication. I probe the former by restricting the sample to applicants whose residence county differs from their processing office county, and the latter using alternative control groups. Section 6.1 presents these tests, and Appendix B.1 provides extended discussion of each threat, including selective attrition, institutional responses, and legal representation.

Third, the immigration office's large-scale hiring during 2015–2017 raises the possibility of differential staffing responses at treated offices, which would conflate behavioral change among incumbent officers with compositional effects of new hires. Without officer-level identifiers in the CRF, I cannot directly separate these channels. Appendix B.1 addresses this concern using register attrition data (Table B.1), finding no evidence of differential attrition across all counties or the subset hosting immigration offices. Section 6.3 provides complementary evidence: the arson effect is concentrated among cases adjudicated through personal interviews, while applicants processed via written procedures at the same office show no comparable response, consistent with a within-officer behavioral channel rather than a purely compositional one.

5.2 Estimation Strategy

The treatment indicator D from the identification section maps to the estimation as $D_{dt} = \mathbf{1}(t \geqslant T_d)$ for TWFE and $D_{dts} = \mathbf{1}(t \geqslant T_{g(s)})$ for the stacking estimator.

I estimate two complementary specifications on a person × month panel: a collapsed difference-in-differences and a dynamic event study. Each is implemented using both a heterogeneity-robust stacking estimator (preferred) following Baker et al. (2022) and a conventional TWFE specification (benchmark). The outcome Y_{irdt} is a binary indicator equal to one if applicant i, residing in county r and processed by branch office d, receives a positive first-instance decision in calendar month t, and zero otherwise. I compare treated offices against not-yet-treated and never-treated offices. Treatment timing varies across the 30 treated offices, creating a staggered adoption design.[5]

Collapsed difference-in-differences. For each treatment cohort g, defined by the calendar month T_g of the first arson attack within 25 km, I construct a sub-experiment pairing the offices treated at T_g with all not-yet-treated and never-treated offices, restricted to the event window $[T_g - 5,\ T_g + 11]$. Stacking the sub-experiments yields:

$$Y_{irdts} = \beta \cdot \mathbf{1}(t \geqslant T_{g(s)}) + \gamma_{d(s)} + \gamma_{rt(s)} + \delta_{t,\tau,n,s} + X_i \Gamma + \varepsilon_{irdts} \tag{5}$$

where s indexes sub-experiments, τ denotes months since application (duration), n denotes nationality, $g(s)$ maps each sub-experiment to its treatment cohort's attack month, and Γ is the coefficient vector on individual controls X_i. The parenthetical (s) on subscripts denotes interaction with the sub-experiment indicator—e.g., $d(s)$ indexes branch office × sub-experiment cells, and $rt(s)$ indexes residence county × calendar month × sub-experiment cells. The fixed effects $\gamma_{d(s)}$ (branch office × sub-experiment), $\gamma_{rt(s)}$ (residence county × calendar month × sub-experiment), and $\delta_{t,\tau,n,s}$ (calendar month × duration × nationality × sub-experiment) are interacted with the sub-experiment indicator, ensuring that each treated cohort is compared only against its own clean control group. X_i includes gender, family status, birth year, and minor status. Standard errors are clustered at the branch office level. The treatment coefficient β estimates the average change in the monthly probability of a positive decision across the twelve post-treatment months at treated offices. The 48 offices provide the identifying variation, with 30 experiencing a first attack within 25 km during the sample period.

Event-study specification. To trace the temporal profile of the treatment effect and check pre-trends, I replace the post-treatment indicator with event-time dummies:

[5] Similar identification strategies have been applied to assess mortality rates instead of approval rates (Miller et al., 2021).

$$Y_{irdts} = \sum_{\substack{k=-5 \\ k \neq -1}}^{11} \beta_k \cdot \mathbf{1}(t - T_{g(s)} = k) + \gamma_{d(s)} + \gamma_{rt(s)} + \delta_{t,\tau,n,s} + X_i \Gamma + \varepsilon_{irdts} \tag{6}$$

The coefficients β_k identify differential approval probabilities, within applications that share the same branch office, residence county × calendar month, and calendar month × time since application × nationality cell, between applicants processed by offices exposed to a nearby arson attack k months prior and those processed by not-yet- or never-treated offices in the same sub-experiment.

TWFE specification (comparison). For comparison, I also estimate the conventional TWFE specifications, which pool all offices and event-times into a single regression without sub-experiment interactions:

$$Y_{irdt} = \beta \cdot \mathbf{1}(t \geqslant T_d) + \gamma_{rt} + \gamma_d + \delta_{t,\tau,n} + X_i \Gamma + \varepsilon_{irdt} \tag{7}$$

$$Y_{irdt} = \sum_{\substack{k=-5 \\ k \neq -1}}^{11} \beta_k \cdot \mathbf{1}(t - T_d = k) + \gamma_{rt} + \gamma_d + \delta_{t,\tau,n} + X_i \Gamma + \varepsilon_{irdt} \tag{8}$$

The person-month panel with a binary outcome and duration dependence connects naturally to discrete-time hazard models. The linear probability specifications in Equations 6 and 8 can be read as linear approximations to such a model, with the duration fixed effects capturing the baseline hazard.

6 Results

The preferred stacking estimator yields a reduction of 2.3 percentage points in the monthly probability of a positive asylum decision at offices exposed to a nearby arson attack (Table 1, Column 3). The estimate is broadly consistent across specifications. The conventional TWFE with and without residence county fixed effects produces estimates of similar magnitude (Columns 1 and 2).[6]

[6] The stacking estimator constructs separate sub-experiments for each treatment cohort, which explains why it has more person-month observations (2,062,257 vs. 918,000) but fewer unique persons (93,950 vs. 148,268) than the TWFE specification: control observations are duplicated across sub-experiments, while persons whose entire panel falls outside all cohort event windows are excluded.

Table 1: Effect of arson attacks on the monthly probability of a positive asylum decision (collapsed post-period)

	(1) Baseline	(2) Residence FE	(3) Stacking
Post × Treated	-0.019***	-0.024***	-0.023***
	(0.007)	(0.006)	(0.007)
Mean Dep. Var.	0.045	0.045	0.046
N (Person × Month)	920,098	918,000	2,062,257
N (Persons)	148,373	148,268	93,950
Residence × month FE	No	Yes	Yes

Notes: Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Each column reports the coefficient on a post-treatment indicator (Equation 5) that collapses the event-study post-period into a single effect. Column (1): baseline TWFE with branch office and calendar month × duration × nationality fixed effects. Column (2): TWFE adding residence county × calendar month fixed effects. Column (3): preferred stacking estimator (Equation 5) following Baker et al. (2022), which constructs clean comparison samples of not-yet-treated and never-treated offices for each treatment cohort to rule out forbidden comparisons under staggered timing. All specifications include demographic controls (gender, family status, birth year, minor status). Standard errors clustered at the branch office level in parentheses.
*p<0.10, **p<0.05, ***p<0.01

Figure 3 presents the dynamic event-study coefficients from the stacking estimator. The pre-treatment coefficients are statistically insignificant. Effects become statistically significant at $k = +2$ and increase through month 4, remaining elevated through month 6. This delayed onset is consistent with the institutional lag between asylum interviews, where adjudication is concentrated, and the recording of decisions (see Appendix Figures A.1 and A.2 for processing durations by stage). The interview of applicants is the moment at which officer discretion is highest and where treatment is most likely to affect it. Arson attacks that precede the interview are most likely to affect the outcome of an application, explaining the delayed onset in approval probabilities. The period in which the effects are highest (months 4–7) mirrors the average delay between interview and the rendering of a decision (see Figure A.2 in the Appendix). The pattern for full refugee status is nearly identical (Appendix Figure C.1).

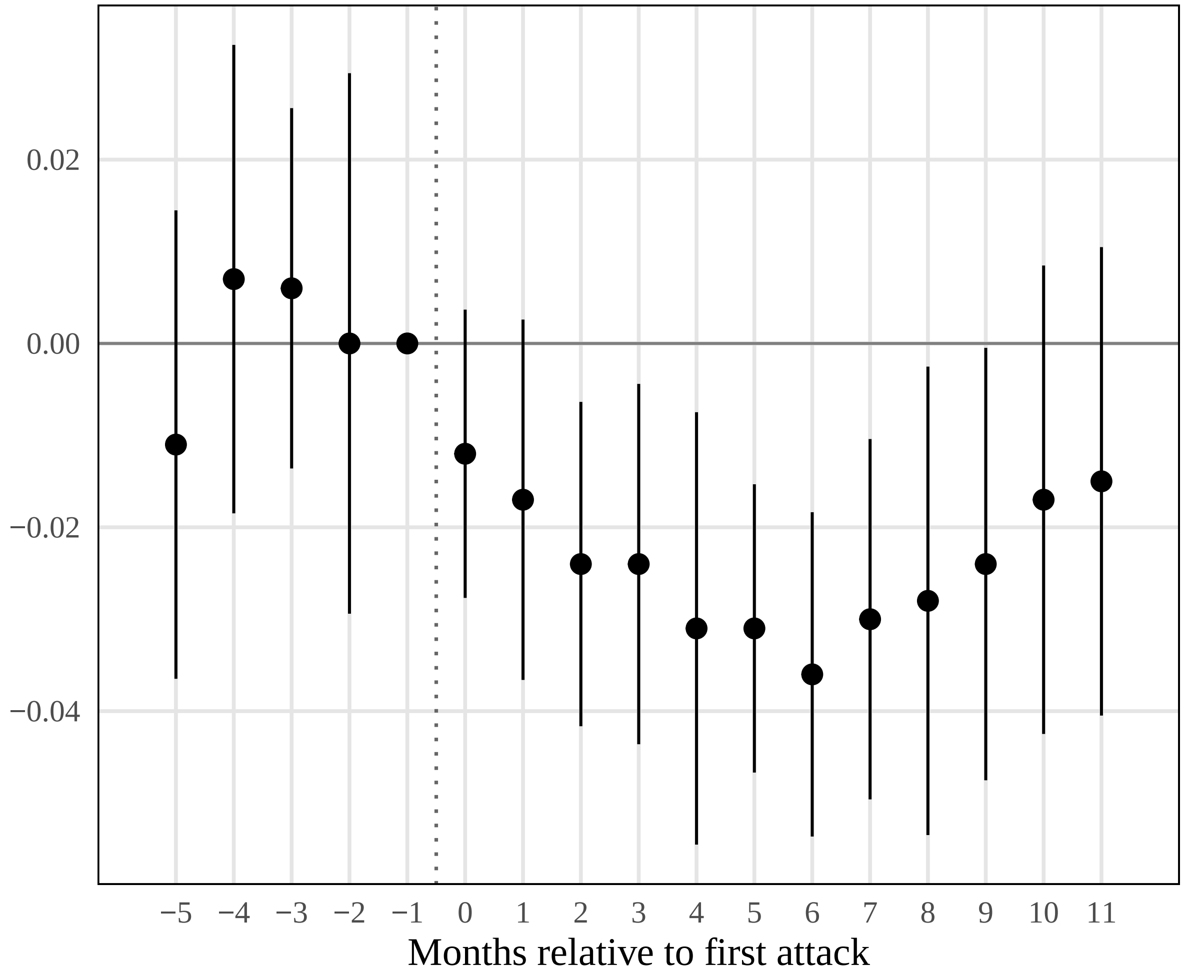

Figure 3: Effect of Arson Attacks on the Monthly Probability of a Positive Asylum Decision
Notes: Stacking estimator (Equation 6) following Baker et al. (2022). Coefficients show change in the probability of receiving any form of protection relative to the month before treatment ($k = -1$). Shaded areas indicate 95% confidence intervals. Standard errors clustered at the branch office level. Sample: 2016–2018. Source: 20% random sample from the CRF; parliamentary interpellations (arson events).

The effect varies across applicant groups (Table C.4). It is largest among applicants from major refugee-sending countries and among nationalities with higher baseline recognition rates, where the monthly hazard has more room to decline. I detect no effect among nationalities with very low recognition rates. Two interpretations are consistent with this pattern. The first is a mechanical floor, where low baseline rates leave little room for decline. The second is differential discretion, where officers exercise genuine judgment primarily in cases where the legal framework does not near-determine the outcome.

The effect is also larger for single women than for single men, consistent with greater scope for discretion in claims that rely on subjective credibility assessments. The effect is concentrated in areas with below-median AfD vote share and absent in high-AfD areas. This could reflect that officers in areas where anti-immigrant sentiment is already high have already adjusted their baseline behavior, leaving no additional margin for the arson signal to shift.

6.1 Robustness

The most direct robustness check restricts the sample to applicants whose residence county differs from their processing office county; the effect persists, reinforcing the interpretation of an office-side channel rather than an applicant-side response to local violence. Table 2 shows that effects are stable across alternative sample timeframes (restricting to decisions through 2016, 2017, or the full window) and distance thresholds (10 km, and 25 km), though narrower radii yield less precise estimates.

Table 2: Robustness: Sample Timeframe and Distance Thresholds

	Timeframe			Distance	
	(1) 2015–16	(2) 2015–17	(3) 2015–18	(4) 10km	(5) 25km
Post × Treated	-0.011* (0.006)	-0.013** (0.006)	-0.014** (0.006)	-0.011 (0.016)	-0.032*** (0.009)
N (Person × Month)	1,010,150	1,276,785	1,357,787	1,356,103	994,688
N (Persons)	149,127	160,259	166,494	171,932	148,768

Notes: Each column reports the mean post-treatment coefficient (event-time periods $k = 0, \ldots, 11$) from the preferred event-study specification (any protection, never-treated comparison). Columns (1)–(3) vary the sample timeframe; columns (4)–(5) vary the distance threshold. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects and demographic controls. *p<0.10, **p<0.05, ***p<0.01.

Extending the observation window to 18 or 24 months leaves point estimates essentially unchanged (Appendix Figure C.3). This addresses the concern that the CRF's recording of only legally binding status changes could conflate suppressed approvals with delayed ones. Wild cluster bootstrap inference on the collapsed difference-in-differences coefficient confirms significance at the 5% level (Appendix C.3). Decomposing the stacking estimator by treatment-timing cohort, the post-treatment effect is negative and statistically significant for all four 2016 cohorts (March–June), with point estimates ranging from −2.2 to −2.8 percentage points (Appendix Figure C.4). The two later cohorts (April and June 2017) yield smaller, statistically insignificant estimates near zero, consistent with the institutional tightening of oversight documented in Section 2.2.

6.1.1 Office-Level Aggregate Statistics

A complementary check on the stacking estimator results uses aggregate statistics on actual asylum decisions observed at the immigration office level.[7] These data record the total number and outcomes of decisions rendered by each office, nationality, and year

[7] I thank Martin Lange and co-authors for sharing the office-level aggregate data from Barreto et al. (2022).

cell. Because the data are aggregated at the office level, they lack the demographic and processing-stage detail of the CRF sample used in the main analysis. Applicant behavior cannot be disentangled from immigration officer behavior at this level of aggregation. I therefore use these data as a robustness check rather than a primary specification.

The office-level data are well suited to verify two assumptions underlying the main identification. The aggregate counts include only first-instance decisions and exclude appeals, so any effect detected here cannot be driven by differential appeal behavior across treated and untreated offices. The data also allow a direct test of Assumption ID.3. If arson attacks shifted the timing of decisions rather than their content, the number of decisions at treated offices should decline. I construct a continuous exposure measure $E_d = (13 - m_d)/12$, where m_d is the calendar month of office d's first attack, so that a January attack yields $E_d = 1$ (full-year exposure) and a December attack yields $E_d \approx 0.08$.

Table 3 reports estimates from the specification

$$y_{dnt} = \alpha_d + \gamma_{nt} + \beta\, E_d + \varepsilon_{dnt}, \tag{9}$$

where y_{dnt} is the outcome for office d, nationality n, year t. α_d are office fixed effects and γ_{nt} are nationality $\times$ year fixed effects. The sample covers 2015 to 2018 and excludes Syrian, Iraqi, and Eritrean nationalities, who received almost universal approval in 2015 through written procedures and therefore lack a pre-treatment period. Treated offices are restricted to pre-treatment and treatment-year observations. Columns (1) and (2) use the protection rate as the outcome. Columns (3) and (4) use log decisions to test whether attacks shifted the volume of decisions.

The exposure coefficient on the protection rate is negative and marginally significant without state $\times$ year controls (column 1) and strengthens with their inclusion (column 2). This confirms the main result in an independent dataset that excludes appeals. The log-decisions placebo in columns (3) and (4) shows no significant effect on decisions at treated offices. If anything, results point to larger volumes of decisions rather than slower processing. The decision-rendering hazard does not appear to have increased at treated offices, consistent with Assumption ID.3. Arson attacks shifted the content of decisions, not their timing or volume.

Table 3: Office-Level Aggregate Statistics

	(1)	(2)	(3)	(4)
	Any protection		Log decisions	
Exposure	−0.043*	−0.068*	0.272	0.120
	(0.025)	(0.039)	(0.270)	(0.345)
Mean dep. var.	0.175	0.175	3.732	3.732
N	2,206	2,206	2,206	2,206
Offices	44	44	44	44
Treated	22	22	22	22
Office FE	Yes	Yes	Yes	Yes
Nationality × year FE	Yes	Yes	Yes	Yes
State × year FE	No	Yes	No	Yes

Notes: Exposure = (13 − attack month)/12, so that earlier attacks yield higher exposure. Sample: 2015–2018, excluding Syrian, Iraqi, and Eritrean nationalities. Year-of-treatment sample: treated offices observed only in pre-treatment and treatment years; never-treated offices observed in all years. Columns (1)–(2) weight by number of decisions; columns (3)–(4) are unweighted. Standard errors clustered at the office level. Data: aggregate office-level decision statistics (first-instance decisions only, excluding appeals). *p<0.10, **p<0.05, ***p<0.01.

6.2 Exploratory Evidence on Channels

I report two exploratory dimensions along which the treatment effect varies: news coverage of the attack and attack severity.

6.2.1 News coverage

To measure the salience of arson attacks, I identify news articles covering each attack using the GDELT Global Knowledge Graph, filtering by arson and refugee-related themes within a one-day window of each incident. I then count the number of distinct national newspapers reporting on the attack, using IVW-audited circulation records (Q4 2014) to identify newspaper outlets. Figure 4 presents interaction coefficients from an event study interacting event-time dummies with this continuous newspaper count. The interaction is negative and significant in months 1–3 after treatment, indicating that attacks covered by more newspapers are associated with a larger reduction in approval rates during the early post-treatment window. The effect fades after month 4 and is absent at longer horizons. This transient pattern is consistent with, but does not identify, an information channel in which media salience amplifies the initial signal of anti-refugee sentiment; alternative interpretations — correlated attack severity, correlated local political climate, or selection into press coverage — cannot be ruled out with the present data.

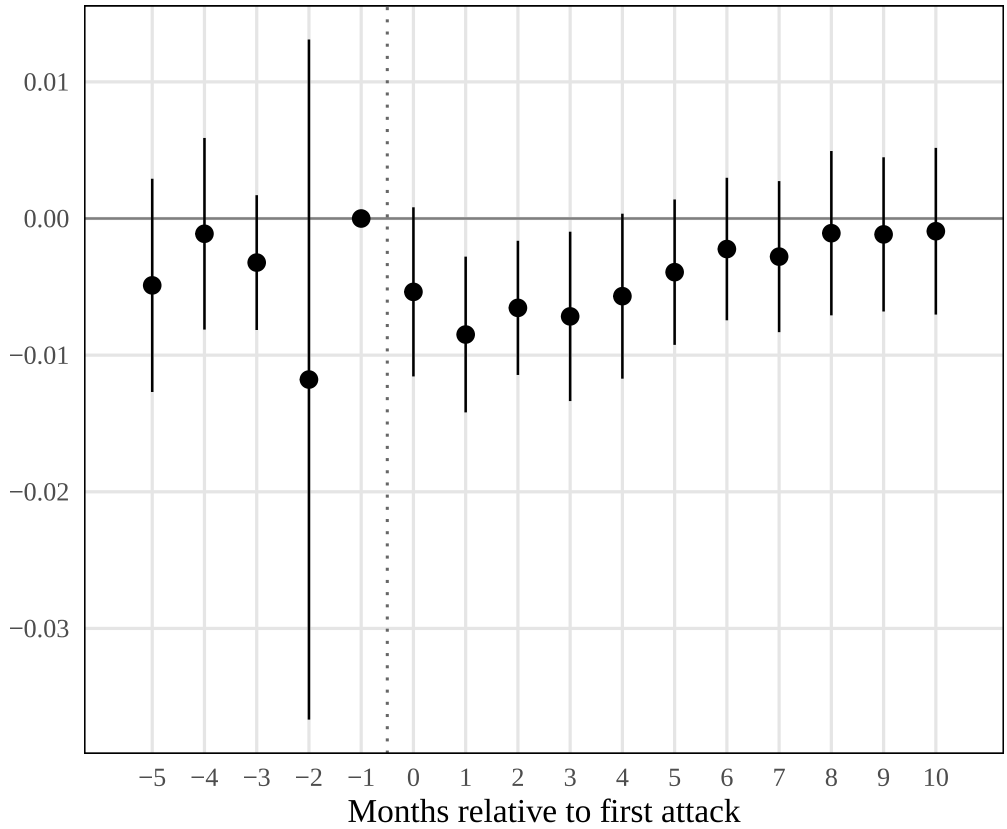

Figure 4: News Coverage Interaction

Notes: Interaction coefficients from event study interacting event-time dummies with continuous newspaper count (number of national newspapers reporting on the attack). Dependent variable: any protection. Reference period: $k = -1$. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects. Standard errors clustered at the branch office level. Shaded area shows 95% confidence intervals. Sample: 2016–2018. Source: 20% random sample from the CRF; GDELT Global Knowledge Graph; parliamentary interpellations (arson events).

6.2.2 Attack severity

A second exploratory cut examines whether the severity of attacks is associated with a different treatment profile. Using the criminal code classification from the interpellation data, I distinguish between simple arson and aggravated arson, which involves greater risk to human life.[8] The aggravated arson interaction shows a contrasting temporal pattern: the interaction coefficient is insignificant in early post-treatment months but turns positive and significant at months 6 and 8–10, a pattern loosely consistent with severe attacks triggering a countervailing response that partially offsets the negative effect at longer horizons. The pattern connects to research showing that extreme tactics can provoke sympathetic countervailing responses (Feinberg et al., 2020; Wasow, 2020); whether the severity distinction operates in the same way in the asylum-adjudication setting is a question the present data cannot settle.

[8] See Appendix Table A.1 for the full classification of criminal offenses.

6.2.3 Public attitudes toward refugees

I use survey data from the SOEP to test how attacks against refugees shifted beliefs near immigration offices. Only counties that host an immigration office branch enter the sample. I construct a stacked difference-in-differences design paralleling the main specification, comparing SOEP respondents in treated immigration office counties against respondents in never-treated immigration office counties. The event-study specification is:

$$\tilde{W}_{ics} = \sum_{\substack{k=-3 \\ k\neq -1}}^{8} \beta_k \cdot \mathbf{1}(\text{month}_s = k) \cdot \text{Treated}_c + \phi \tilde{W}_{i,t-1} + \alpha_{c(s)} + \mu_{k(s)} + \varepsilon_{ics} \tag{10}$$

where $\tilde{W}_{ics}$ is the worry outcome for respondent i in county c and sub-experiment s, standardized by the pre-treatment standard deviation. Because interview dates vary across respondents within each survey wave, the SOEP functions as a repeated cross-section at the respondent level for this exercise: each respondent contributes one observation per wave, and identification comes from comparing respondents interviewed before and after the attack within the same county. Treated_c indicates whether the respondent resides in a county whose immigration office was exposed to a nearby arson attack; all respondents in the sample reside in counties that host an immigration office branch. $\tilde{W}_{i,t-1}$ is the lagged outcome from the previous survey wave. The fixed effects $\alpha_{c(s)}$ (cohort × treated) and $\mu_{k(s)}$ (cohort × month bin) are interacted with the sub-experiment indicator. Standard errors are clustered at the county level.

Figure 5 presents event-study coefficients for respondents residing in treated immigration office counties, expressed in pre-treatment standard deviations. Pre-treatment coefficients are statistically insignificant for both outcomes. If anything, immigration worry displays a declining trend before treatment, making the post-attack increase a conservative estimate. The two outcomes display distinct temporal patterns at onset. Xenophobia worry spikes by roughly 0.3 SD in the month of the attack but drops back to near zero by month 1 and remains flat thereafter. Immigration worry rises by about 0.15 SD at impact and stays elevated over months 0–3 before reaching 0.20–0.25 SD by months 4–5, though confidence intervals grow substantially at longer horizons. Arson attacks trigger an immediate concern about xenophobic hostility that fades as the event recedes from public attention, while worries about immigration accumulate as the local salience of the refugee issue persists.

The short-lived increase in concern for refugee safety aligns with Freitas-Monteiro and Prömel (2024), who show a short-term increase in pro-refugee beliefs after large far-right protests at the national level, and with Emeriau (2024), who documents a short-run empathy response in approval rates after migrant shipwrecks. The more persistent pattern of immigration worry is also consistent with the voting responses to far-right terrorist attacks documented by Sabet et al. (2023) and with the longer-lasting decrease in approval rates found in the main results.

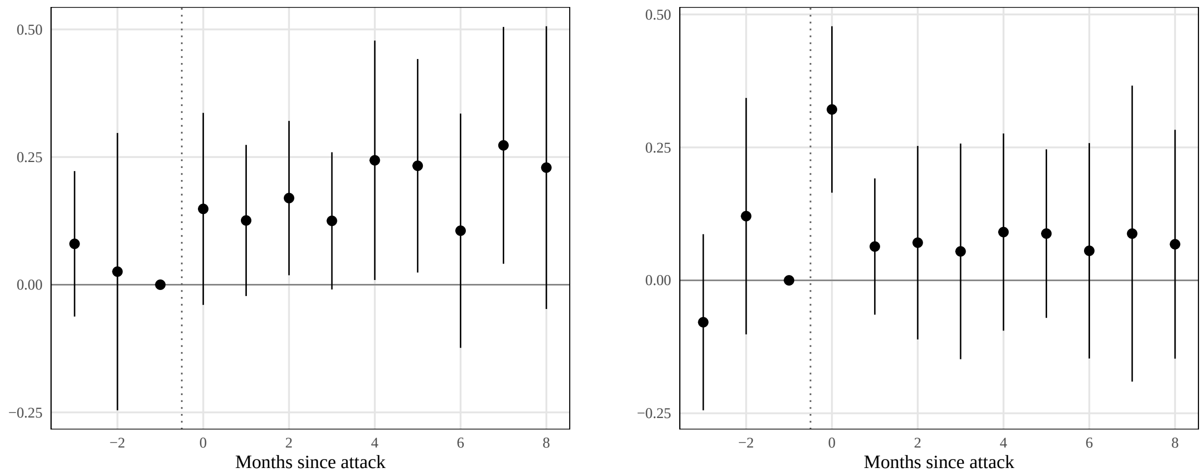

(a) Immigration worry (b) Xenophobia worry

Figure 5: Event study: public attitudes in treated counties
Notes: Stacked difference-in-differences event-study coefficients. Dependent variables are SOEP worry items (1–3 scale, increasing in concern), standardized by the pre-treatment standard deviation. Treatment: first arson attack in the respondent's county of residence. Control group: respondents in never-treated immigration office counties. Reference period: month −1. All specifications include cohort×treated and cohort×month-bin fixed effects, controlling for lagged immigration worry, xenophobia worry, and social cohesion worry. Standard errors clustered at the county level. Error bars show 95% confidence intervals. Data: SOEP v41, 2015–2019.

6.3 Personal Interview as the Site of Discretion

The main results rely on the institutional argument that officer discretion is concentrated at the personal interview. A separate policy change provides direct evidence for this claim. During 2015, the immigration office processed claims from Syrian and Iraqi-minority applicants through simplified written procedures without individual interviews. Beginning in 2016, personal interviews were reinstated: applicants who arrived after 1 January 2016 were required to undergo a personal hearing, immigration office headquarters later extended this requirement to all persons filing from 17 March 2016, regardless of entry date.[9] This creates within-office variation in interview format among applicants processed at the same branch office. The specification includes branch-office-by-nationality, calendar-month, and duration-by-time-by-nationality fixed effects, along with personal-interview-by-branch-office fixed effects that absorb the average level difference in outcomes between interview and non-interview applicants within each office. Identification of the interaction therefore relies on whether the arson effect differs across interview formats within the same office. The *Asylpaket II* bundled interview reinstatement with other policy changes (safe-country designations, family reunification restrictions); the fixed effects absorb any uniform effects of the broader policy package, isolating the interview-format

[9]The strict treatment definition excludes Iraqi applicants, who were only partially subject to the original written procedure; the broad definition includes them.

margin.

Table 4 interacts the post-arson-attack indicator with the personal interview indicator, testing directly whether the arson effect operates through the interview stage. Under both treatment definitions, the arson effect on acceptance rates is concentrated among applicants whose claims were adjudicated through a personal interview; applicants processed via the written procedure at the same office during the same period show no comparable response to nearby attacks. This confirms that the interview is the channel through which local hostility enters adjudication: officers who never meet applicants face-to-face do not adjust their decisions in response to arson exposure, while those conducting personal hearings do.

Table 4: Arson Effect by Interview Format

	(1) Strict	(2) Broad
Post-Attack	−0.0018	−0.0009
	(0.0030)	(0.0033)
Post-Attack × Personal Interview	−0.0136**	−0.0114***
	(0.0054)	(0.0039)
Observations	440,192	440,192
N individuals	75,595	75,595

Notes: Interaction of the post-arson-attack indicator with a personal interview indicator. Dependent variable: positive first-instance decision. "Personal Interview" indicates applicants subject to the personal hearing requirement. Column (1) uses the strict definition excluding Iraqi applicants only partially subject to the original written procedure; column (2) uses the broad definition including them. Sample restricted to the three major refugee-sending nationalities (Syria, Afghanistan, Iraq) with positive baseline acceptance rates. Controls include gender, family status, birth year, and minor status. Standard errors clustered at the individual level. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

7 Conclusion

This paper uses geographic separation between processing offices and applicant residences in the German asylum system to test whether far-right violence compromises high-stakes administrative decisions. Arson attacks against refugee accommodations reduce the monthly approval probability at nearby branch offices, and the effect is concentrated among nationalities where officers have genuine latitude over outcomes. A within-office test confirms the personal hearing as the site of discretion: caseworkers who meet applicants face-to-face adjust their decisions after nearby attacks, while those deciding cases through written procedures at the same office do not. The tightening of internal quality assurance after 2017 coincides with attenuated effects in later treatment cohorts, consistent with institutional oversight constraining this response. Far-right violence can, in sum, shift the outcomes of a process designed to reflect the merits of each individual claim.

What makes these findings distinct is that the hostility signal is *in-group* violence against a vulnerable out-group. Unlike out-group-perpetrated terrorism (Shayo and Zussman, 2011; Brodeur and Wright, 2019), arson attacks on refugee housing offer no rational basis for updating threat assessments, yet the bureaucratic response mirrors them. Survey evidence indicates that compliance dominates resistance as responses to political violence. The hostile signal accumulates in the local environment while sympathy for victims fades, providing a plausible attitudinal pathway through which the political climate enters the adjudication room. The results point toward concrete safeguards. Structured interviews with standardized question catalogues could limit the scope for subjective assessment. Systematic training and feedback of the kind Emeriau (2023) shows can reduce bias among asylum officers would address the individual-level channel, while rotation of caseworkers across regions could insulate offices from persistent local conditions. More broadly, wherever public servants exercise discretion within the boundaries of formal rules (Lipsky, 2010), their decisions may absorb the political climate of their environment—a vulnerability that extends well beyond immigration policy.

References

Aksoy, C. G., P. Poutvaara, and F. Schikora (2023, September). First time around: Local conditions and multi-dimensional integration of refugees. *Journal of Urban Economics 137*, 103588. 2.1

Arnold, D., W. Dobbie, and C. S. Yang (2018, November). Racial Bias in Bail Decisions. *The Quarterly Journal of Economics 133*(4), 1885–1932. 1

Baker, A. C., D. F. Larcker, and C. C. Y. Wang (2022, May). How much should we trust staggered difference-in-differences estimates? *Journal of Financial Economics 144*(2), 370–395. 5.2, 1, 3

BAMF-Forschungsdatenzentrum (2023). AZR Research Dataset 2021. 12

Barreto, C., P. Berbée, K. Gallegos Torres, M. Lange, and K. Sommerfeld (2022, April). The Civic Engagement and Social Integration of Refugees in Germany. *Nonprofit Policy Forum 13*(2), 161–174. 7

Bolhaar, J., N. Ketel, and B. Van Der Klaauw (2020, March). Caseworker's discretion and the effectiveness of welfare-to-work programs. *Journal of Public Economics 183*, 104080. 1

Brodeur, A. and T. Wright (2019, December). Terrorism, immigration and asylum approval. *Journal of Economic Behavior & Organization 168*, 119–131. 1, 7

Bursztyn, L., G. Egorov, and S. Fiorin (2020, November). From Extreme to Mainstream: The Erosion of Social Norms. *American Economic Review 110*(11), 3522–3548. 1

Chen, D. L. and M. Loecher (2025, June). Mood and the malleability of moral reasoning: the impact of irrelevant factors on judicial decisions. *Journal of Behavioral and Experimental Economics 116*, 102364. 1

Chen, D. L., T. J. Moskowitz, and K. Shue (2016, August). Decision Making Under the Gambler's Fallacy: Evidence from Asylum Judges, Loan Officers, and Baseball Umpires*. *The Quarterly Journal of Economics 131*(3), 1181–1242. 1

Chenoweth, E. and M. J. Stephan (2011). *Why Civil Resistance Works: The Strategic Logic of Nonviolent Conflict.* Columbia University Press. Google-Books-ID: kGW8cops3GcC. 1, 3

Danziger, S., J. Levav, and L. Avnaim-Pesso (2011, April). Extraneous factors in judicial decisions. *Proceedings of the National Academy of Sciences 108*(17), 6889–6892. 1

Dobbie, W., J. Goldin, and C. S. Yang (2018, February). The Effects of Pre-Trial Detention on Conviction, Future Crime, and Employment: Evidence from Randomly Assigned Judges. *American Economic Review 108*(2), 201–240. 1

Emeriau, M. (2023, October). Learning to be Unbiased: Evidence from the French Asylum Office. *American Journal of Political Science 67*(4), 1117–1133. 1, 3, 7

Emeriau, M. (2024, October). Victim or Threat? Shipwrecks, Terrorist Attacks, and Asylum Decisions in France. *American Journal of Political Science 68*(4), 1187–1204. 1, 2.3, 3, 6.2.3

Eren, O. and N. Mocan (2018, July). Emotional Judges and Unlucky Juveniles. *American Economic Journal: Applied Economics 10*(3), 171–205. 1

Feinberg, M., R. Willer, and C. Kovacheff (2020). The activist's dilemma: Extreme protest actions reduce popular support for social movements. *Journal of Personality and Social Psychology 119*(5), 1086–1111. 1, 3, 6.2.2

Freitas-Monteiro, T. and C. Prömel (2024, December). Local far-right demonstrations and nationwide public attitudes towards migration. *Journal of Economic Behavior & Organization 228*, 106766. 1, 3, 6.2.3

Gleiser, P. and C. Salas Escobar (2024). AZR-Forschungsdatensatz 2023. 4.2, 12

Gould, E. D. and E. F. Klor (2016, November). The Long-run Effect of 9/11: Terrorism, Backlash, and the Assimilation of Muslim Immigrants in the West. *The Economic Journal 126*(597), 2064–2114. 1

Grosjean, P., F. Masera, and H. Yousaf (2023, February). Inflammatory Political Campaigns and Racial Bias in Policing*. *The Quarterly Journal of Economics 138*(1), 413–463. 1, 3

Gundacker, L., Y. Kosyakova, and G. Schneider (2024, December). How regional attitudes towards immigration shape the chance to obtain asylum: Evidence from Germany. *Migration Studies 13*(1), mnae002. 1

Hemker, J. and A. Rink (2017, October). Multiple Dimensions of Bureaucratic Discrimination: Evidence from German Welfare Offices. *American Journal of Political Science 61*(4), 786–803. 1

Heyes, A. and S. Saberian (2019, April). Temperature and Decisions: Evidence from 207,000 Court Cases. *American Economic Journal: Applied Economics 11*(2), 238–265. 1

Jaschke, P., S. Sardoschau, and M. Tabellini. NBER WORKING PAPER SERIES. 2.1

Jomni, H. and N. Zakharov (2026, January). Do terrorist attacks polarize politicians? Evidence from the European parliamentary speeches on migration. *Journal of Peace Research 63*(1), 86–97. 1

Lipsky, M. (2010). *Street-Level Bureaucracy, 30th Ann. Ed.: Dilemmas of the Individual in Public Service*. Russell Sage Foundation. 1, 3, 7

MacKinnon, J. G., M. . Nielsen, and M. D. Webb (2023, February). Cluster-robust inference: A guide to empirical practice. *Journal of Econometrics 232*(2), 272–299. C.3

Maestas, N., K. J. Mullen, and A. Strand (2013, August). Does Disability Insurance Receipt Discourage Work? Using Examiner Assignment to Estimate Causal Effects of SSDI Receipt. *American Economic Review 103*(5), 1797–1829. 1

Miller, S., N. Johnson, and L. R. Wherry (2021, August). Medicaid and Mortality: New Evidence From Linked Survey and Administrative Data*. *The Quarterly Journal of Economics 136*(3), 1783–1829. 5

Roth, J. (2022). Pretest with caution: Event-study estimates after testing for parallel trends. *American Economic Review: Insights 4*(3), 305–322. 5.1

Sabet, N., M. Liebald, and G. Friebel (2023). Terrorism and Voting: The Rise of Right-Wing Populism in Germany. *SSRN Electronic Journal*. 1, 3, 6.2.3

Shayo, M. and A. Zussman (2011, August). Judicial Ingroup Bias in the Shadow of Terrorism *. *The Quarterly Journal of Economics 126*(3), 1447–1484. 1, 7

Spirig, J. (2023, January). When Issue Salience Affects Adjudication: Evidence from Swiss Asylum Appeal Decisions. *American Journal of Political Science 67*(1), 55–70. 1, 3

Wasow, O. (2020, August). Agenda Seeding: How 1960s Black Protests Moved Elites, Public Opinion and Voting. *American Political Science Review 114*(3), 638–659. 3, 6.2.2

A Institutional Details and Data

A.1 Institutional Details

Internal Quality Assurance. The immigration office operates a dual-review principle under which a second reviewer checks each asylum decision before dispatch. During the 2015–2016 caseload surge, internal oversight was substantially weakened: the central quality-assurance unit reviewed approximately 1% of the 282,700 decisions rendered in 2015, and the dual-review requirement was made formally binding by central directive only in September 2017.[10] A mandatory random 10% pre-dispatch quality check was introduced in May 2018 following the discovery of procedural irregularities at the Bremen branch office. These dates imply that for the first half of the sample period (2016–mid-2017), the institutional constraints on individual officer discretion were weaker than the standing procedural framework would suggest.

[10] Source: parliamentary inquiries BT-Drs. 19/4853 and testimony by the immigration office staff council before the Interior Committee in 2018.

Table A.1: Classification of Arson Offenses (German Criminal Code)

Category	StGB §	German Title	Official English Title	Key Statutory Elements
Simple arson				
Simple	§ 306	Brandstiftung	Arson	Sets fire to or destroys buildings, plants, warehouses, motor vehicles, forests, or agricultural facilities belonging to another (1–10 years)
Simple	§ 306d	Fahrlässige Brandstiftung	Negligent arson	Acts negligently in cases under § 306(1) or § 306a(1), or causes danger under § 306a(2) by negligence ($\leqslant$5 years)
Aggravated arson				
Aggravated	§ 306a	Schwere Brandstiftung	Aggravated arson	Sets fire to premises serving to accommodate people, churches, or premises where people are usually present ($\geqslant$1 year)
Aggravated	§ 306b	Besonders schwere Brandstiftung	Especially aggravated arson	Causes serious health damage or places another in danger of death; acts to facilitate another offence or prevents extinguishing ($\geqslant$2 years; $\geqslant$5 years for danger of death)
Aggravated	§ 308	Herbeiführen einer Sprengstoffexplosion	Causing explosion	Causes an explosion, in particular using explosives, and thereby endangers life, limb, or property of significant value ($\geqslant$1 year)

Notes: Classification of arson offenses under the German Criminal Code (*Strafgesetzbuch*, StGB). Section titles and statutory elements are from the official English translation provided by the Federal Ministry of Justice (gesetze-im-internet.de). Penalty ranges in parentheses. Simple arson covers § 306 and § 306d. Aggravated arson covers § 306a, § 306b, and § 308 — the categories that endanger human life, used as the severity indicator in the mechanism analysis.

Table A.2: Types of International Protection in Germany

Protection Status	**Requirements**	**Legal Basis**	**Residence Duration**	**Family Reunification**	**Permanent Residency**
Asylum (Art. 16a GG)	Political persecution by state actors; must enter directly from a non-safe country	German Basic Law (Grundgesetz)	3 years	Yes (easier, no waiting period)	After 3 years
Refugee Protection	Persecution due to race, religion, nationality, political opinion, or social group	Geneva Refugee Convention, §3 AsylG	3 years	Yes (easier, no waiting period)	After 3 years
Subsidiary Protection	No individual persecution but serious harm (e.g., death penalty, torture, war-related threats)	§4 AsylG	1 year (extendable to 3)	Restricted (2-year waiting period, proof of financial means required)	After 5 years
National Ban on Deportation	Risk of inhumane treatment or life-threatening situation in home country, but no refugee/subsidiary protection status	§60 Abs. 5/7 AufenthG	1 year (renewable)	Restricted (must prove financial stability)	After 7 years

Notes: Overview of protection categories under German asylum law as of 2016–2018 (the sample period). Legal bases: German Basic Law (Art. 16a GG), Asylum Act (AsylG), and Residence Act (AufenthG). Residence permit durations and family reunification rules reflect the legal framework during the study period. Source: Federal Office for Migration and Refugees (BAMF) procedural guidelines.

A.2 Descriptive Statistics

Table A.3: Descriptive Statistics: Asylum Applications

	Person-months	Unique persons
Panel A: Sample construction		
Raw panel	5,057,154	210,715
After duration restriction (⩽12 mo. of filing)	2,528,568	210,714
After sample-period restriction (2016–2018)	2,077,182	207,438
After event window [−5, +10] or never-treated	1,187,918	159,213
Estimation sample: at-risk for first pos. decision	920,341	148,405
Panel B: Treatment assignment in estimation sample		
Never-treated (comparison, all periods)	449,957	55,758
Not-yet-treated (comparison, pre-treatment)	46,254	16,420
Treated (post first arson within 25 km)	407,955	85,717
Panel C: Branch offices		
Ever treated within sample window	30	
Never treated within sample window	18	
Total branch offices	48	

Notes: Construction of the estimation sample from the 20% CRF random sample. Unit of observation is the applicant-month. Source: 20% random sample from the CRF; immigration office branch-office register; parliamentary interpellations (arson events).

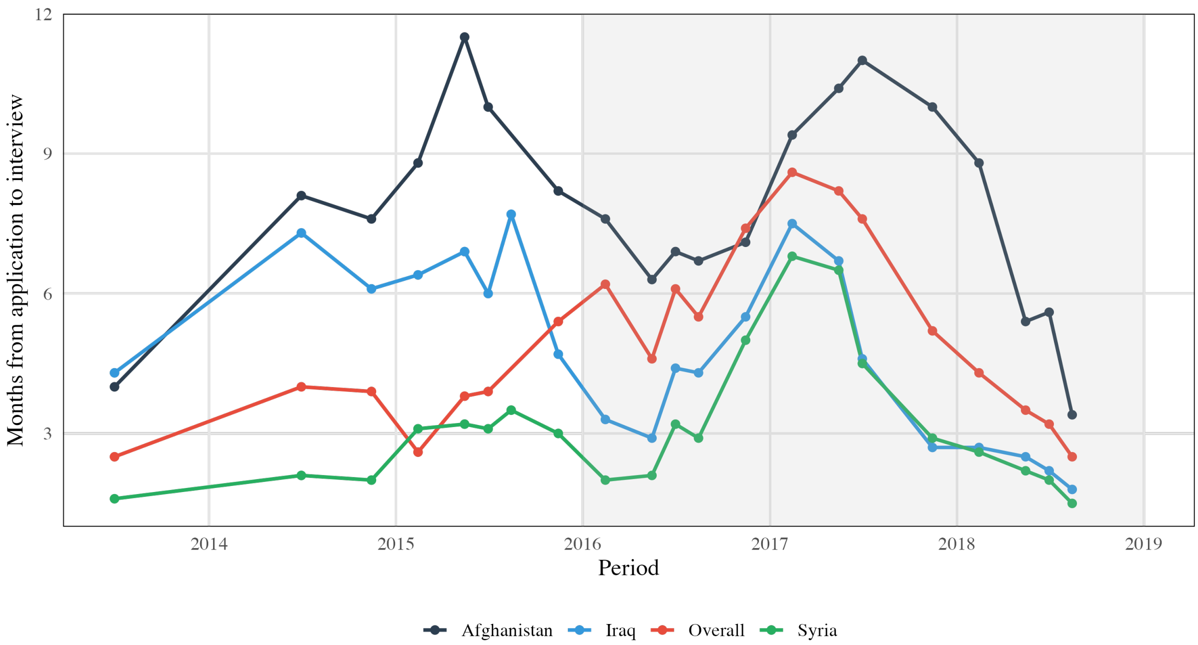

Figure A.1: Duration from Application to Interview

Notes: Mean months from asylum application filing to personal interview (*Anhörung*), by quarter. Lines show quarterly averages for the overall caseload and for Syria. Shaded region marks the analysis sample (2016–2018). Gaps indicate quarters not covered by the parliamentary inquiry source. Source: Federal parliamentary inquiry data on immigration office processing times, 2013–2018.

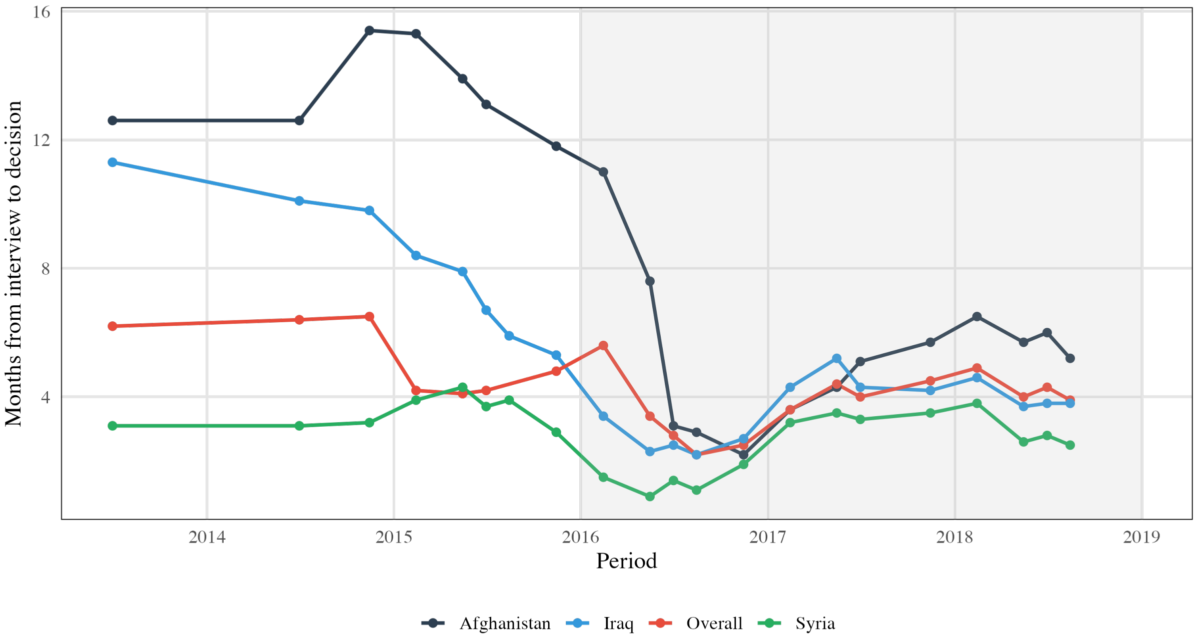

Figure A.2: Duration from Interview to Decision

Notes: Mean months from personal interview (*Anhörung*) to final branch-office decision, by quarter. Lines show quarterly averages for the overall caseload and for Syria. Shaded region marks the analysis sample (2016–2018). Gaps indicate quarters not covered by the parliamentary inquiry source. Source: Federal parliamentary inquiry data on immigration office processing times, 2013–2018.

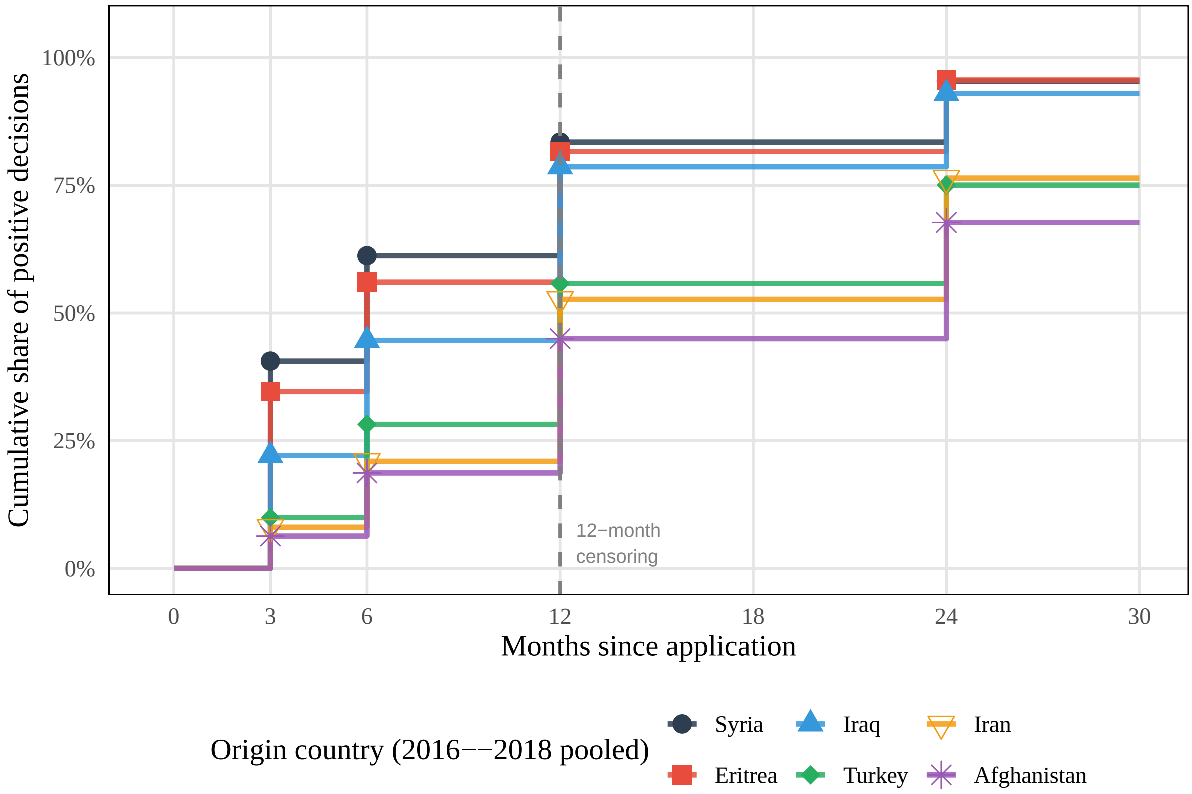

Figure A.3: Cumulative Distribution of Time to Positive Decision by Origin Country

Notes: Cumulative share of positive asylum decisions rendered within a given number of months after application filing, by major origin country. Data pooled over application years 2016–2018. Each step function traces one nationality; markers indicate the observed thresholds (3, 6, 12, and 24 months). The dashed line marks the 12-month censoring horizon used in the main estimation sample. Syria and Eritrea are resolved fastest (over 80% of positive decisions within 12 months); Afghanistan and Iran are slowest (under 50% within 12 months). Nigeria is omitted due to missing cell counts. Source: 20% random sample from the CRF.

B Identification

B.1 Identification Threats

This subsection provides detailed discussion of each identification concern introduced in Section 5.1.

Allocation of refugees. The identification strategy leverages the quasi-random allocation of refugees across German counties. While certain counties specialize in processing applications from specific countries of origin, these specializations remain largely stable over time. The key identifying assumption is that the distribution of refugees across German regions does not systematically change in correlation with local anti-immigrant sentiment. Even if counties attempted to select applicants based on local attitudes, the specifications control for all refugee characteristics available for selection, including country of origin, gender, age, and family status. Figure B.1 provides direct evidence on this point: predicted approval rates, computed from cell means of nationality, calendar month, gender, and minor status, show no shift around treatment timing.

Selective attrition. The CRF does not retain records of individuals who naturalize. If naturalization rates differ between treated and control counties, the 2023 extract would disproportionately lose approved cases from the group with higher naturalization, biasing observed approval rates downward for that group.[11] The structural timeline limits this concern in either direction. The standard naturalization path requires 6-8 years of lawful residence, so applicants recognized in 2016 would satisfy the residence requirement in 2022 at the earliest. The results are robust to controlling for place of residence, meaning that for this channel to drive the results, it would need to operate through the processing location rather than residence. Table B.1 tests this directly by comparing treated and control counties on two administrative margins: register attrition, defined as the log ratio of protection-title holder counts between the 2021 and 2023 CRF extracts,[12] and cumulative naturalizations as a share of the county's foreign population. Neither measure differs significantly between treated and control counties, whether considering all counties or restricting to those hosting an immigration office. The latter comparison more closely mirrors the unit of treatment variation in the main analysis.

[11] For attrition to generate the estimated negative treatment effect, treatment would need to *increase* naturalization among affected counties—removing more approved cases from the treated group and depressing its observed approval rate. The reverse channel, in which treatment lowers approval rates and thereby reduces naturalization eligibility, would attenuate the estimate toward zero.

[12] The main analysis uses the 2023 CRF extract (Gleiser and Salas Escobar, 2024). An earlier 2021 extract (BAMF-Forschungsdatenzentrum, 2023), drawn before any applicant in the 2016–2018 cohort could have satisfied the residence requirement for naturalization, provides a pre-naturalization benchmark. This extract contains fewer variables and could not be used for the primary analysis, but suffices for comparing holder counts across cohorts. The ratio of holder counts between the two extracts captures all sources of register attrition.

Exogeneity of attack timing. The identifying assumption requires that the timing of first arson attacks near immigration offices is uncorrelated with unobserved determinants of approval rates, conditional on the fixed effects. Arson attacks are not randomly assigned: they cluster in areas with rising anti-immigrant sentiment. The conceptual framework treats arson attacks as salient signals through which officers learn about local anti-immigrant sentiment. The identification concern is not that sentiment exists, but that offices may already be responding to the same underlying sentiment through other channels, such as colleague conversations, local media tone, or community attitudes, before the attack occurs. If so, the attack would be a symptom of sentiment that has already influenced decisions rather than an independent informational shock. Three features of the design mitigate this concern. First, the calendar month × nationality × duration fixed effects absorb nationwide sentiment shifts. Second, branch office fixed effects absorb time-invariant local attitudes. Third, the residence county × calendar month fixed effects absorb both cross-county heterogeneity and time-varying shocks at the county level in the comparison group. What remains unabsorbed is time-varying sentiment *specific to the branch office's locality* that correlates with attack timing. I cannot fully rule out this threat. The event study's pre-treatment coefficients provide indirect evidence: in the preferred specification, none of the lead coefficients is statistically significant (Figure 3).

Spillover. The event study design assumes that arson attacks near one branch office do not affect decisions at other offices. Two types of spillover are relevant. First, geographic spillover: news coverage of an attack may affect applicants' behavior or legal-support availability near the attack site, independent of the processing office. The residence county × calendar month fixed effects absorb time-varying shocks at applicants' locations, and the restriction to applicants whose residence county differs from their processing office county directly probes this channel. Second, office-to-office spillover: staff at untreated offices may respond to attacks at distant offices through internal immigration office communication or national media coverage. If such spillover depresses approval rates at control offices, the treatment effect is attenuated; if it elevates them (e.g., through a compensatory response), the estimate is inflated. Under non-uniform spillover, the dynamic coefficients can be distorted in either direction, and attenuation is not guaranteed. Results are robust to alternative control groups—never-treated offices only, and the full not-yet-treated pool—which would be differentially affected if internal spillover were pervasive. Nevertheless, spillover biases cannot fully be ruled out.

Institutional responses. Immigration office leadership may have responded to arson attacks with internal directives, changes in oversight intensity, or personnel reallocation—channels distinct from individual officer discretion. If such institutional responses systematically lowered approval rates at affected offices, the estimated effect would conflate individual behavioral change with top-down institutional reactions. As discussed in Section 5.1, the federal structure of asylum adjudication limits any branch-office-level managerial reaction to procedural adjustments rather than substantive policy change. Without access to internal communications or personnel records, however, I cannot rule out procedural responses such as shifting case assignments or intensifying supervisory review at affected offices. Following the discovery of procedural irregularities at the Bremen branch office in mid-2018, the immigration office introduced systematic monitoring of branch-office approval rates against a country-composition-adjusted reference rate, with deviations exceeding 10 percentage points triggering central review.[13] This reform falls at the trailing edge of the sample period and does not drive the main results—the through-2017 robustness specification predates it entirely—but its introduction confirms that headquarters recognized cross-office variation in adjudication as a quality concern.

Legal representation. Arson attacks may trigger increased legal aid or NGO mobilization near affected areas, which could independently affect approval rates by improving case preparation. Since applicants reside in different counties from their processing offices, this channel would need to operate through the applicant's residence county rather than the office's locality. The residence county × calendar month fixed effects in the preferred specification absorb both time-invariant and time-varying differences in legal support across applicant locations, limiting this channel to within-county-month variation in legal mobilization.

B.2 Balance Tests

The balance test in Figure B.1 checks whether applicant composition, as summarized by predicted approval rates, shifts systematically around the timing of a first nearby arson attack. Predicted approval rates are computed from cell means (nationality × calendar month × gender × minor status) using either pre-treatment periods or the never-treated sample, which summarize an applicant's expected approval probability based on observable characteristics alone. I then regress each predicted rate on the same event-study specification used for the main outcome (Equation 6), restricting to the first observation per application. If treatment status is uncorrelated with applicant composition, the event-time coefficients should be statistically indistinguishable from zero throughout the event window. Significant pre- or post-treatment movements in predicted approval rates would indicate compositional change that the calendar month × duration × nationality fixed effects in the main specification are absorbing. Both variants—using pre-treatment and never-treated cell means—show no evidence of compositional change.

[13]The reference rate weights each branch office's nationality mix by the national approval rate per nationality, producing a counterfactual mean against which the office's actual rate is compared (source: parliamentary inquiries BT-Drs. 19/6786 and 20/14882).

Table B.1: Balance: treated vs. control counties

	Treated	Control	Diff.	p
Panel A: All counties				
Register attrition ($\ln N^{2023}/N^{2021}$)	-0.070	-0.029	-0.041	0.261
	(0.022)	(0.029)	(0.036)	
Naturalizations 2020–2023 (% of foreign pop.)	5.68	6.08	-0.40	0.120
	(0.14)	(0.21)	(0.25)	
N	227	140		
Panel B: Immigration office counties				
Register attrition ($\ln N^{2023}/N^{2021}$)	0.046	0.105	-0.059	0.649
	(0.083)	(0.099)	(0.129)	
Naturalizations 2020–2023 (% of foreign pop.)	5.91	5.74	0.17	0.820
	(0.35)	(0.63)	(0.72)	
N	33	16		

Notes: Treated counties experienced at least one arson attack within 25 km. Control counties are never-treated. Register attrition defined as $\ln(N_c^{2023}/N_c^{2021})$ where N_c^t is the total number of protection-title holders in county c across the 2015–2018 recognition cohorts in the CRF extract of year t. Naturalizations are cumulative 2020–2023 counts as a percentage of the county's 2020 foreign population (Destatis Einbürgerungsstatistik). Standard errors of the mean in parentheses. Immigration office counties are those hosting a branch office of the Federal Office for Migration and Refugees. Sample restricted to counties observed in both datasets.

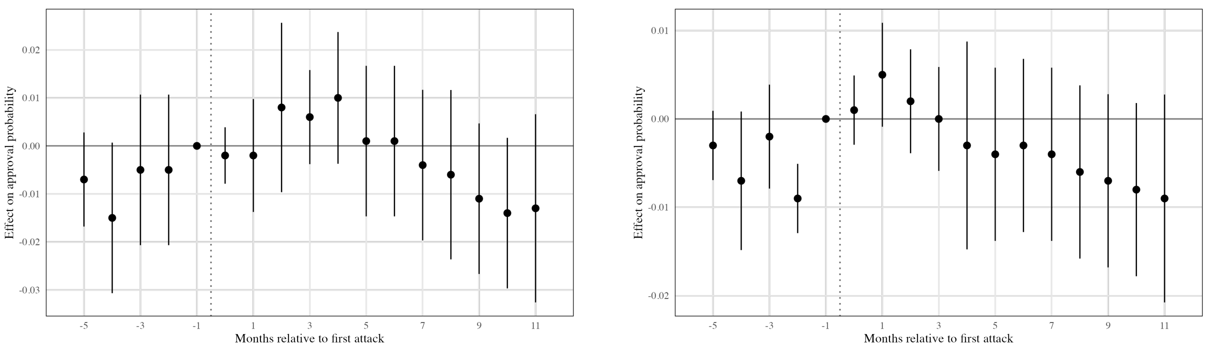

(a) Pre-treatment cell means **(b) Never-treated cell means**

Figure B.1: Balance Tests: Predicted Acceptance Rates around Treatment

Notes: Event-study coefficients from regressions of predicted acceptance rates on event-time dummies, relative to the month before treatment ($k = -1$). Predicted rates are computed from cell means (nationality × calendar month × gender × minor status) using pre-treatment periods (a) or the never-treated sample (b). Sample restricted to the first observation per application. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects. Standard errors clustered at the branch office level. Vertical bars show 95% confidence intervals. Source: 20% random sample from the CRF; parliamentary interpellations (arson events).

B.3 Exposure Construction

This subsection documents the procedure by which branch offices are matched to arson events in the interpellation data. The same procedure is used throughout Section 5 and its robustness variants.

(i) Branch-office geocoding. Each of the 48 immigration office branch offices is represented by a single geographic point. Where the street address of the branch is available in the immigration office register, the point is the address geocoded via the BKZ–AGS crosswalk maintained in the project's data-preparation pipeline; the BKZ (*Behördenkennziffer*) is the administrative identifier for the branch and the AGS (*Amtlicher Gemeindeschlüssel*) is the official municipality key. Where the street address is unavailable or the address geocode is missing, the branch is assigned to the centroid of its hosting municipality (AGS).

(ii) Attack-record geocoding. Arson events drawn from the parliamentary interpellation record list the event's city and the date. The combined AAS / ZEIT-ONLINE "Hass im Netz" compiled dataset, which cross-references the same police records, is used to resolve city names to AGS codes; each event is then placed at the centroid of its AGS. 23 event records are ambiguous in the city-name field (typographic variants, disbanded municipalities, or inter-municipal event locations) and are resolved manually. The list of 23 manual overrides is documented in the replication package and flagged so that Phase 2 analyses can substitute alternative geocodes if needed.

(iii) Tie-break rule. An office's first-treatment event is the earliest qualifying attack within the 25 km radius. When multiple qualifying attacks occur in the same calendar month, the earlier date is chosen; within a single calendar day, the attack with the smaller distance to the branch-office point is chosen. This tie-break is applied only to establish T_d in Equation 8; subsequent attacks fold into the post-treatment regime as discussed in Section 5.1.

(iv) Distance metric. Distance between the branch-office point and the attack point is computed using the Vincenty ellipsoid formula on WGS84 coordinates, which is appropriate for the distance range relevant here (a few hundred metres to several tens of kilometres). The 25 km radius, the 10 km robustness threshold, and the same-county threshold are all computed from this metric.

(v) Office-by-office first-treatment table. A separate table lists each of the 48 branch offices with its first-treatment calendar month under the 10 km, 25 km, and same-county rules.

(vi) Minimum-attack-distance distribution. A companion figure shows, for each branch office, the distribution of minimum distances to the nearest arson event during 2015–2021. The figure identifies, for each branch, which radius threshold the branch lies on the border of — a useful diagnostic for the sensitivity of the treatment classification to small changes in the distance cutoff.

C Additional Results

C.1 Robustness

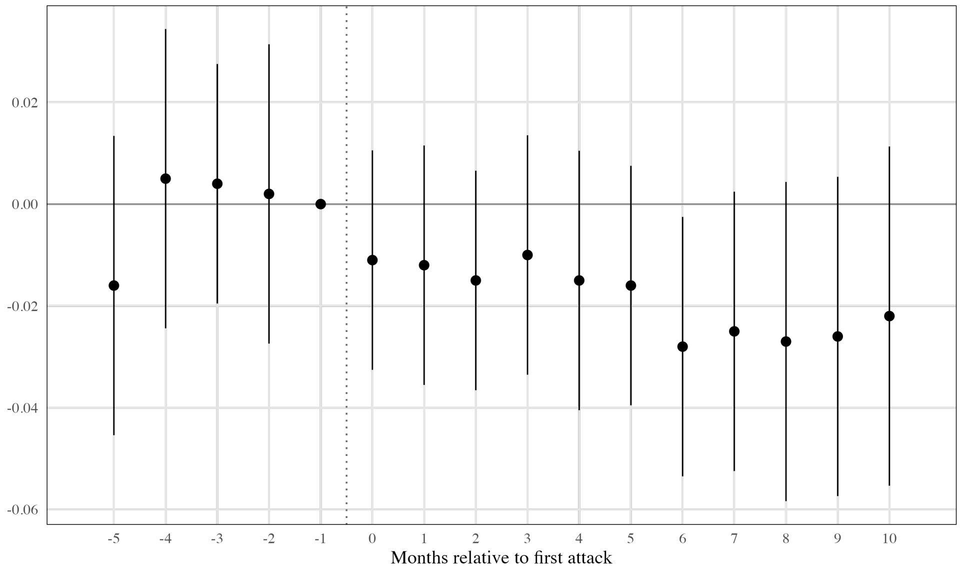

Figure C.1: Event Study: Full Refugee Status

Notes: Coefficients show change in the monthly probability of receiving full refugee status relative to the month before treatment ($k = -1$). Vertical bars indicate 95% confidence intervals. Standard errors clustered at the branch office level. Sample: 148,405 applicants observed over 920,341 person-months in the risk-set window, 2016–2018. See Figure 3 for the corresponding any-protection result. Source: 20% random sample from the CRF; parliamentary interpellations (arson events).

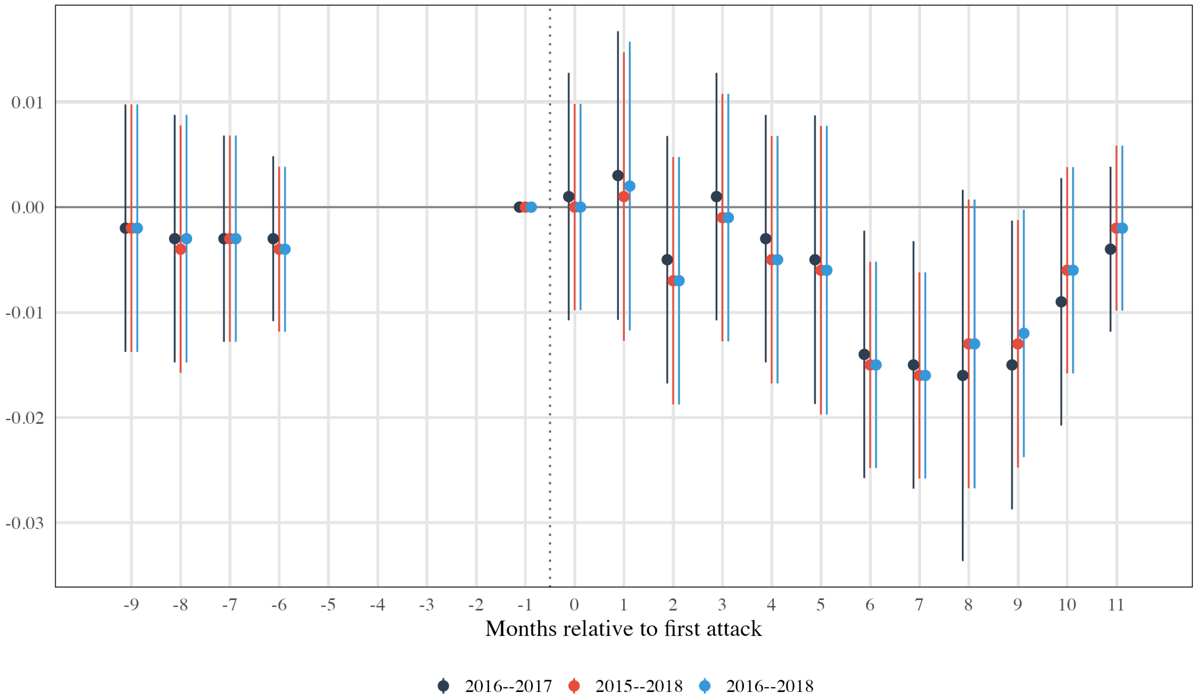

Figure C.2: Robustness: Event study coefficients across sample timeframes (any protection)
Notes: Event-study coefficients for any protection across different sample timeframes. Each series restricts the sample to decisions through the indicated year. Dependent variable: binary indicator for receiving any form of protection. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects, and demographic controls. Standard errors clustered at the branch office level. Vertical bars show 95% confidence intervals. Source: 20% random sample from the CRF; parliamentary interpellations (arson events).

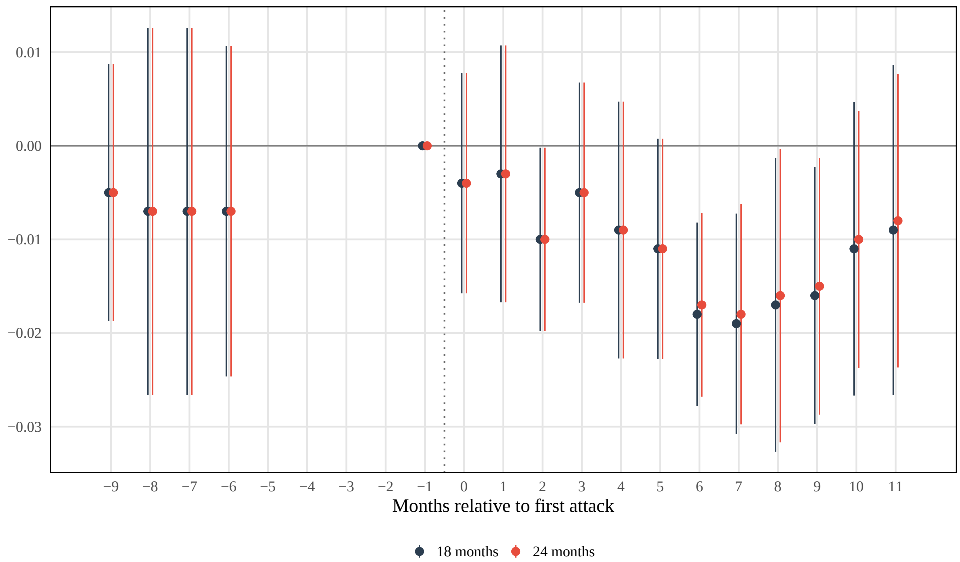

Figure C.3: Robustness: Maximum observation duration
Notes: Event-study coefficients varying the maximum observation duration after application: 12, 18, and 24 months. Dependent variable: any protection. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects. Standard errors clustered at the branch office level. Shaded areas show 95% confidence intervals. Source: 20% random sample from the CRF; parliamentary interpellations (arson events).

C.2 Heterogeneous Effects

Table C.1: Arson attacks and the monthly probability of a positive asylum decision: Collapsed difference-in-differences estimates by subgroup

	Dep. Var.: Any Approval				
	Post Effect	(SE)	N	Persons	Mean
Arson Intensity					
No Aggravated Arson	-0.020**	(0.009)	572,871	79,235	0.043
Aggravated Arson	-0.026***	(0.007)	735,259	117,039	0.047
Far-Right Vote Share					
Below Median	-0.019***	(0.007)	920,098	148,373	0.045
Above Median	0.000	(.)	449,710	55,730	0.044
Applicant Gender					
Not Single Female	-0.018**	(0.007)	844,120	135,610	0.044
Single Female	-0.024***	(0.007)	75,150	12,673	0.051
Approval Probability					
Higher Prob.	-0.035***	(0.013)	549,998	92,244	0.072
Lower Prob.	-0.001*	(0.001)	370,096	56,125	0.004
Nationality					
Non-Majority	-0.003*	(0.002)	601,075	88,470	0.015
Majority SYR/AFG	-0.051**	(0.021)	319,022	59,902	0.100
Continent					
Africa	-0.004	(0.005)	149,974	21,135	0.024
Europe	-0.000	(.)	847	115	0.006
Asia	-0.030**	(0.012)	604,843	97,582	0.059

Notes: Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Each row reports the coefficient on a post-treatment indicator from a separate regression estimated on the indicated subsample. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects, as well as demographic controls (gender, family status, birth year, minor status). Standard errors clustered at the branch office level in parentheses. Standard errors marked (.) indicate insufficient within-group variation for estimation.
*p<0.10, **p<0.05, ***p<0.01

Table C.2: Heterogeneous effects: Demographic splits

	Major Origin		Syria		Single Female		Single Male	
	(1) Non-Major	(2) Major	(3) Non-Syr.	(4) Syrian	(5) Other	(6) Single F	(7) Other	(8) Single M
Post × Treated	-0.003*	-0.051**	[TBD]	[TBD]	-0.018**	-0.024***	-0.022***	-0.014**
	(0.002)	(0.021)	([TBD])	([TBD])	(0.007)	(0.007)	(0.008)	(0.006)
Mean Dep. Var.	0.015	0.100	[TBD]	[TBD]	0.044	0.051	0.052	0.034
N (Person × Month)	601,075	319,022	[TBD]	[TBD]	844,120	75,150	561,330	358,397
N (Persons)	88,470	59,902	[TBD]	[TBD]	135,610	12,673	94,186	54,147

Notes: Standard errors in parentheses. The table reports post-treatment difference-in-differences estimates for subsamples defined by demographic characteristics. Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Major origin countries include applicants from major refugee-sending countries (Iraq, Syria, Afghanistan). Single Female/Male identifies unmarried applicants by gender. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects, and demographic controls (gender, family status, birth year, minor status). Standard errors clustered at the branch office level.

*p<0.10, **p<0.05, ***p<0.01

Table C.3: Heterogeneous effects: Geographic and political splits

	East Germany		Far-Right '13		Far-Right '17		Approval Rate		Aggravated Arson		News Coverage	
	(1) West	(2) East	(3) Low	(4) High	(5) Low	(6) High	(7) High	(8) Low	(9) Light	(10) Heavy	(11) High	(12) Low
Post × Treated	[TBD]	[TBD]	-0.019***	0.000	-0.019***	0.000	-0.055**	-0.001*	-0.020**	-0.026***	[TBD]	[TBD]
	([TBD])	([TBD])	(0.007)	—	(0.007)	—	(0.021)	(0.001)	(0.009)	(0.007)	([TBD])	([TBD])
Mean Dep. Var.	[TBD]	[TBD]	0.045	0.044	0.045	0.044	0.119	0.004	0.043	0.047	[TBD]	[TBD]
N (Person × Month)	[TBD]	[TBD]	920,098	449,710	920,098	449,710	253,777	370,096	572,871	735,259	[TBD]	[TBD]
N (Persons)	[TBD]	[TBD]	148,373	55,730	148,373	55,730	51,727	56,125	79,235	117,039	[TBD]	[TBD]

Notes: Standard errors in parentheses. The table reports post-treatment difference-in-differences estimates for subsamples defined by geographic and political characteristics. Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Far-Right '13/'17 splits by median AfD vote share in federal elections. Approval Rate splits by median nationality-level approval rate. Aggravated Arson indicates attacks classified under § 306a, § 306b, or § 308 StGB. Standard errors marked (—) indicate insufficient within-group variation for estimation. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects, and demographic controls. Standard errors clustered at the branch office level.
*p<0.10, **p<0.05, ***p<0.01

Table C.4: Heterogeneous effects: Post-period DiD summary

Subgroup	N	DiD	SE	p-value
Origin				
Major origin countries (IRQ/SYR/AFG)	319,022	-0.051	(0.021)	0.015
Demographics				
Single female	75,150	-0.024	(0.007)	0.001
Single male	358,397	-0.014	(0.006)	0.020
Political context				
High AfD vote share	449,710	0.000	—	—
Low AfD vote share	920,098	-0.019	(0.007)	0.007
Baseline approval				
Low approval probability	370,096	-0.001	(0.001)	0.317
High approval probability	253,777	-0.055	(0.021)	0.009
Attack severity				
Aggravated arson	735,259	-0.026	(0.007)	0.000
Non-aggravated arson	572,871	-0.020	(0.009)	0.026

Notes: Summary of post-period difference-in-differences estimates across heterogeneity dimensions. Dependent variable: binary indicator equal to one if the applicant receives any form of protection in a given month. Each row corresponds to a subgroup estimated from a separate regression with a single post-treatment indicator. Columns report sample size (N), DiD estimate, standard error, and p-value. Entries of 0.000 with no standard error indicate insufficient within-group variation for estimation. All specifications include branch office, residence county × calendar month, and calendar month × duration × nationality fixed effects, and demographic controls. Standard errors clustered at the branch office level. Exact p-values reported.

C.3 Staggered-DiD Diagnostics

Stacking Estimator. This subsection reports the event-study coefficients from the stacking estimator described in Section 5.2 (Equation 6). For each treated cohort, the sub-experiment pairs the treated offices with not-yet-treated and never-treated offices over the same event-time window; the cohort-specific estimates are pooled with inverse-variance weights. Figure 3 plots these coefficients for the preferred specification (any protection, never-treated comparison).

Cohort-Level Effects. To assess whether the aggregate treatment effect is driven by a single cohort or is broadly shared, I estimate the post-treatment DiD coefficient separately for each treatment-timing cohort. Figure C.4 plots the cohort-specific ATT against treatment timing. The four 2016 cohorts (March through June) produce negative and statistically significant effects between −2.2 and −2.8 percentage points. The April 2017 cohort yields a positive but imprecisely estimated coefficient, and the June 2017 cohort is near zero. The attenuation for later cohorts is consistent with the institutional reforms described in Section 2.2: the immigration office strengthened internal oversight during 2017, plausibly reducing the scope for local sentiment to influence decisions.

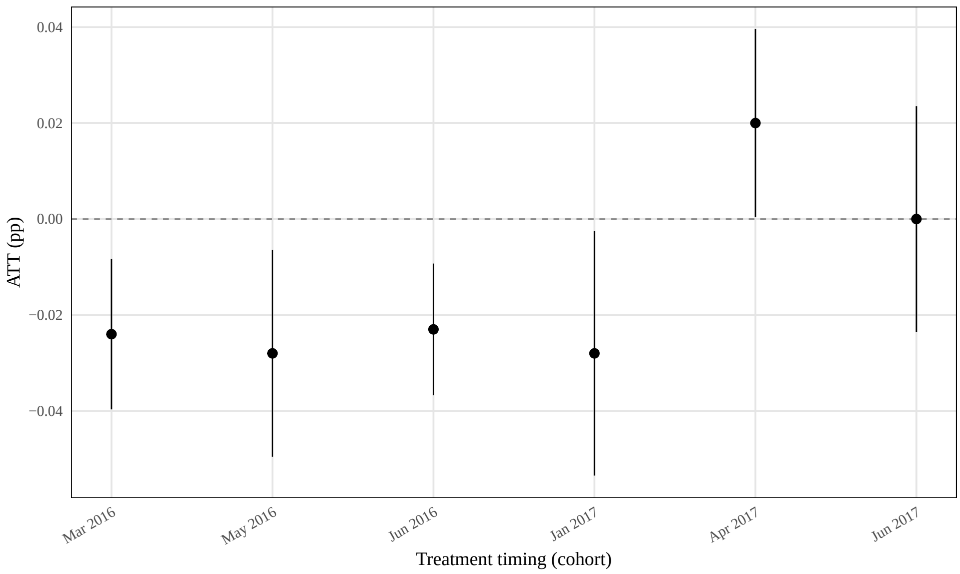

Figure C.4: Cohort-level ATT from stacking estimator

Notes: Each point reports the post-treatment DiD coefficient from a cohort-specific sub-experiment of the stacking estimator (Equation 6). The comparison group consists of never-treated offices. Vertical bars show 95% confidence intervals based on standard errors clustered at the branch office level. The January 2016 cohort is excluded because it contains a single treated office, yielding an imprecise estimate.

Wild Cluster Bootstrap. With 30 treated clusters, conventional cluster-robust standard errors may over-reject the null of zero effect (MacKinnon et al., 2023). I re-compute inference for the collapsed difference-in-differences coefficient — the post-treatment indicator from a regression of residualized approval rates on the treatment dummy — using a wild cluster bootstrap-t procedure with 1,000 Rademacher-weight replications (MacKinnon et al., 2023). The bootstrap p-value is 0.046 ($t(47) = -2.49$), and the 95% confidence set is $[-0.031, -0.0002]$, confirming that the aggregate treatment effect remains statistically significant at the 5% level after small-sample correction.